

# Image Processing

Javier González Jiménez

## Reference Books:

- *Computer Vision: Algorithms and Applications*. Richard Szeliski. Springer. 2010.  
<http://szeliski.org/Book>
- *Visión por Computador*. Javier Gonzalez Jimenez. Thomson-Paraninfo, 1999.

# Content

1. Introduction
2. IP tools and concepts
  - Image color
  - Image Histogram
  - Look-up-tables
  - Distance between pixels
  - Convolution
3. Image Smoothing
4. Image Enhancement

# 1. Introduction


**Goal:** Improve the quality of the image by attenuating noise, adjusting colors, modifying contrast and brightness, ...

In CV this may be a necessary step to **prepare the image for a better feature extraction or segmentation**

An image can be seen as:



a two-dimensional matrix  $A[x,y]$

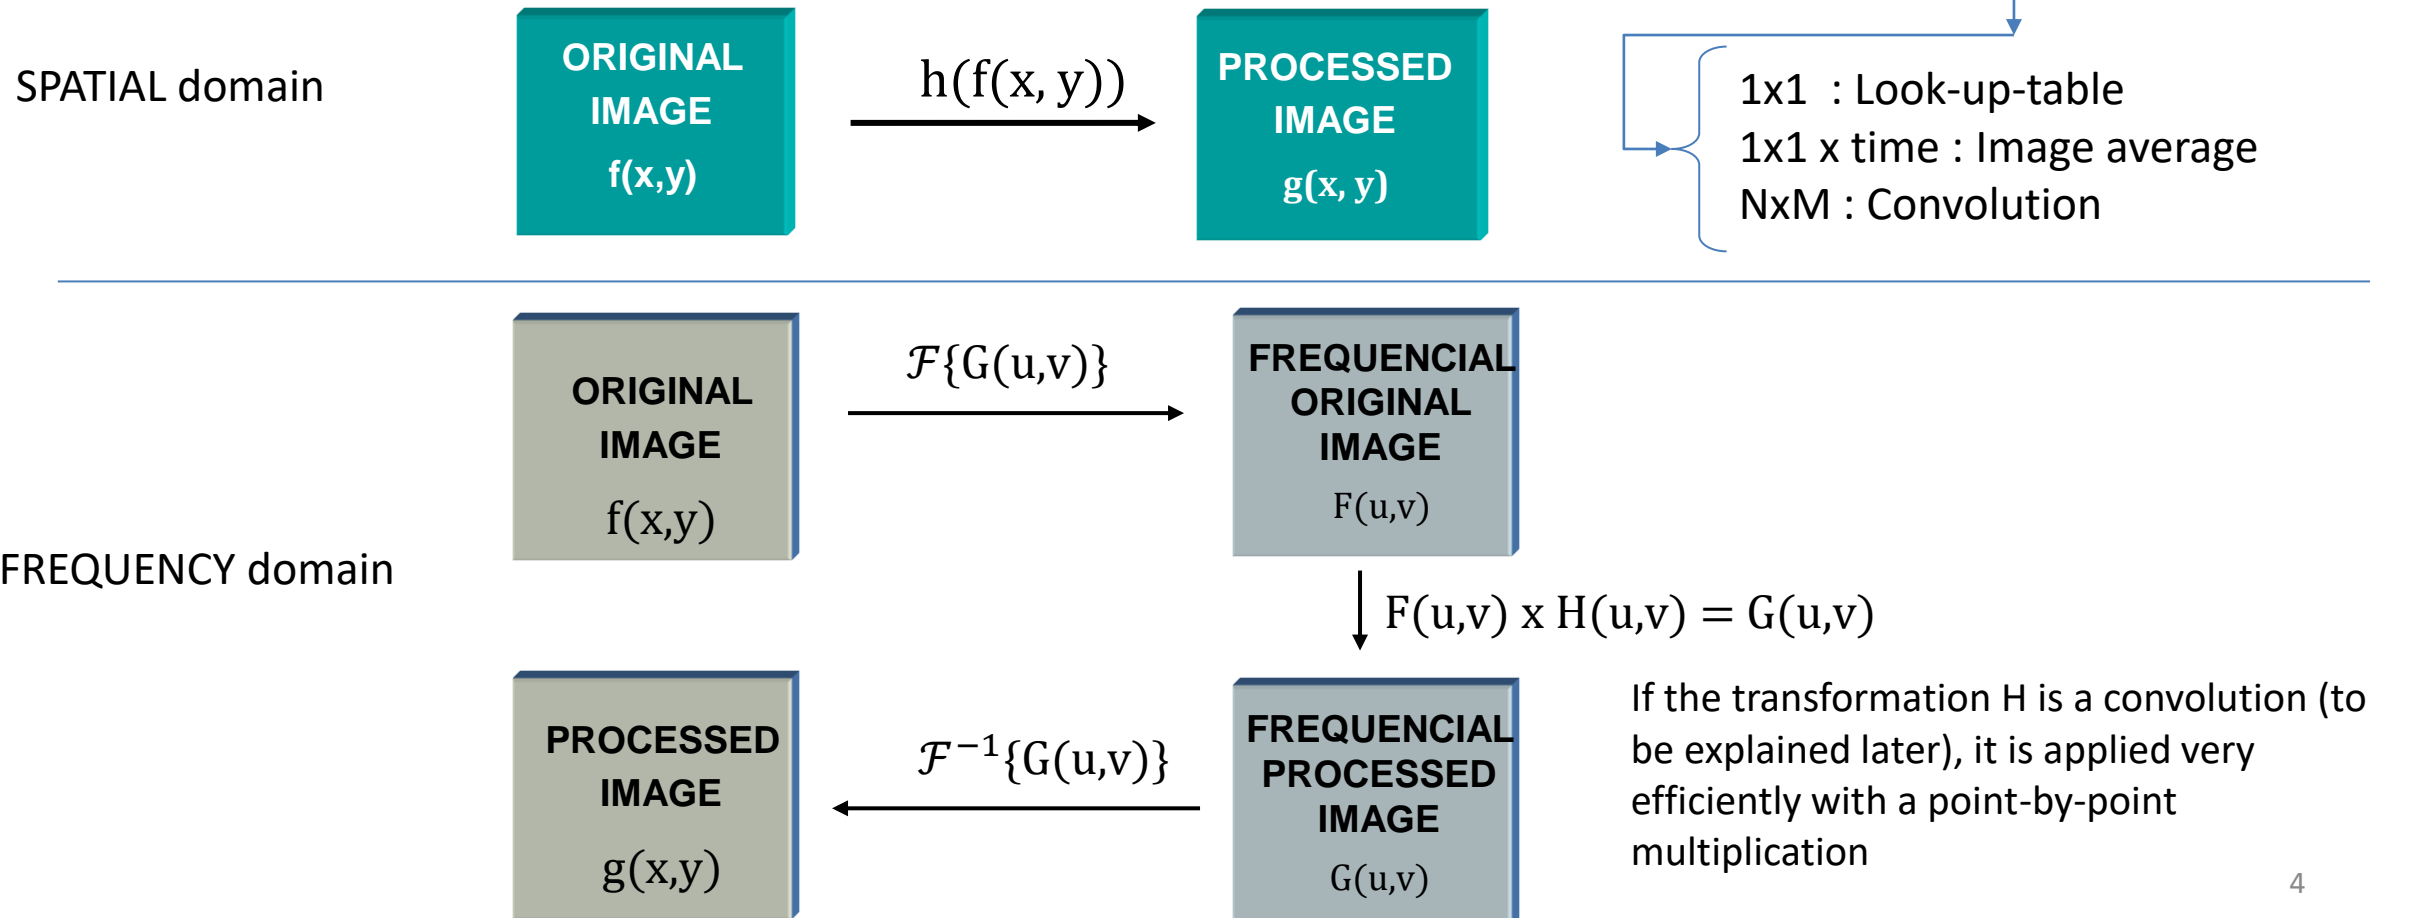


a two-dimensional function  $f(x,y)$  with  $(x,y)$  discrete

Both are the same thing!

# 1. Introduction


Smoothing and enhancement can be applied in the SPATIAL and in the FREQUENCY domain




## 2. IP tools and concepts

Image color

Electromagnetic spectrum



All these wavelengths are produced by some light source (sun or artificial)



Human Luminance Sensitivity Function





We see these wavelengths because the receptors (cones) in our eyes

# Image color

## Human eye



- Two types of photoreceptors in the human retina: **Rods** responsible for intensity, **cones** responsible for color
- **Fovea** - Small region (1 or 2°) at the center of the visual field containing the highest density of cones (and no rods).

Rods and cones are *non-uniformly* distributed on the retina (Less visual acuity in the periphery)




# Image color space

RGB: Linear color space used by computers



Any color can be computed as the linear combination of the three components RGB





When a color has the same value of the three components RGB (diagonal of the RGB trihedron) we have grey-levels.

$R=G=B \rightarrow$  grey level (or gray scale)

# Image color

A color image



Mathematically, a color image is a tri-dimensional array:  $I(\text{row}, \text{column}, \text{color})$  with  $\text{color} \in \{R, G, B\}$

Easiest way to convert from RGB color to grey-level (grayscale):

**Average method:** Grayscale =  $(R + G + B / 3) \rightarrow 33\% R, 33\% G, 33\% B$




## Image color: Chromaticity

A color (R,G,B) can be converted to normalized colors ( $r,g,b$ ) which gives us the proportion of red, green and blue:

$$r = \frac{R}{R+G+B} \quad g = \frac{G}{R+G+B} \quad b = \frac{B}{R+G+B} \quad r + g + b = 1$$

One constraint

Chromaticity chart is a 2D space, for example,  $\langle r, g \rangle$



In the chromaticity chart, the missing component ( $b$ ) is the value of the pixel.


For example: each 45 deg. line (e.g.  $r+g = 0.8 \rightarrow b = 0.2$ ).

This representation is useful to segment regions based on their color

## 2. IP tools and concepts

### Image Histogram

- Is a **representation of the frequency** each color intensity appears in the image
- Built by **counting the occurrence of each color** in the image
- A **color image** has 3 histograms (e.g. R, G, B)
- Provides **statistical information** of the intensity distribution (e.g. brightness and contrast)



## Brightness and contrast (moments of a histogram)

Moment  $k$  of a function  $f(x)$ :  $m_k = \int_{-\infty}^{\infty} x^k f(x) dx$



Central moment  $k$  of a function  $f(x)$ :  $\mu_k = \int_{-\infty}^{\infty} (x - m_1)^k f(x) dx$

**Brightness:** Average of the pixel intensities = One order moment of the histogram ( $h(i)$ )

$$m_1 = \bar{i} = \sum_{i=0}^{255} i^1 h(i)$$

**Contrast:** Standard deviation of the pixel intensities. Square root of the second central moment of the histogram

$$\sigma = \sqrt{\sum_{i=0}^{255} (i - \bar{i})^2 h(i)}$$

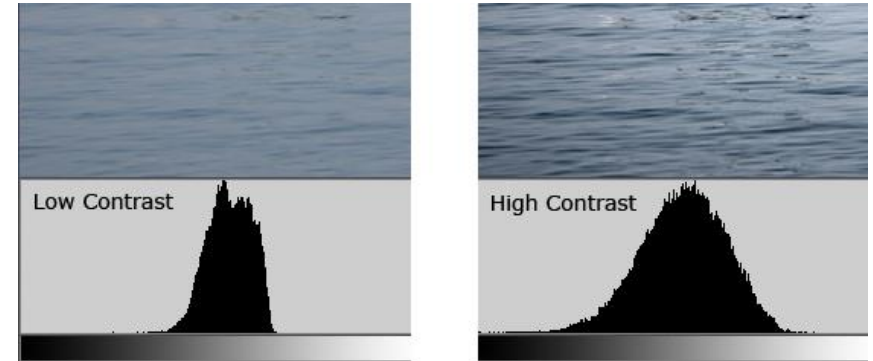


# Image histogram is useful for:


- Detecting Black or White saturation




- Contrast and brightness



- Use of the whole intensity range (i.e. the 256 values)

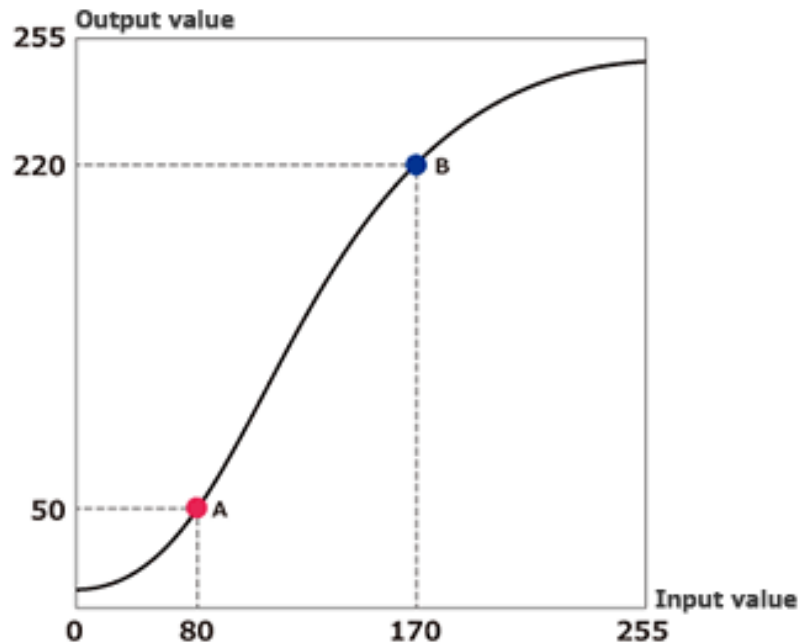


- Select a suitable threshold for image binarization



# Look-up Tables (LUTs)

A LUT replaces an index intensity (input) by another intensity value stored in the table (vector)



Continuous LUT

A table titled 'Discrete LUT (a vector)' showing a mapping of input values to output values. The table has two columns: 'Input value' and 'Output value'. The first 80 rows are not shown, indicated by a bracket and the text '256 values'. The row for input 80 is highlighted in red, and the row for input 170 is highlighted in blue. Point A is marked next to the red row, and point B is marked next to the blue row.

Input value	Output value
0	5
1	6
2	8
3	10
...	...
79	47
80	50
81	53
...	...
169	218
170	220
171	222
...	...
252	248
253	250
254	252
255	254

Discrete LUT (a vector)

All the pixels with intensities 80 in the image are replaced by the value 50

**Example:** Processed image



- does not change the geometry of objects in the image
- only changes the histogram of the image

# Look-up Tables (LUTs)

## RGB LUT

LUT for R		LUT for G		LUT for B	
Input value	Output value	Input value	Output value	Input value	Output value
0	5	0	10	0	0
1	6	1	12	1	0
2	8	2	13	2	1
3	10	3	15	3	1
79	47	79	72	79	25
80	50	80	75	80	25
81	53	81	77	81	26
252	242	252	251	252	234
253	245	253	252	253	236
254	248	254	253	254	238
255	251	255	255	255	240

256 values

3 sets



Applied to an image it is like replacing the original color palette by another one










# Look-up Tables (LUTs)

## Examples







Original image



# Look-up Tables (LUTs)

## Implementation (example binarization)






# Distance between two pixels $p_1, p_2$ :

- $D(p_1, p_2)$  is a distance function if:
1.  $D(p_1, p_2) \geq 0$  [ $D(p_1, p_2) = 0$  si  $p_1 = p_2$ ]
  2.  $D(p_1, p_2) = D(p_2, p_1)$
  3.  $D(p_1, p_3) \leq D(p_1, p_2) + D(p_2, p_3)$


- *Euclidean distance ( $L^2$ -norm)* between  $p_1$  y  $p_2$  :

$$D_e(p_1, p_2) = \left[ (x_1 - x_2)^2 + (y_1 - y_2)^2 \right]^{1/2}$$



- *Manhattan distance ( $l_1$ -Norm)* between  $p_1$  y  $p_2$  :


$$D_4(p_1, p_2) = |x_1 - x_2| + |y_1 - y_2|$$



2
2 1 2
2 1 0 1 2
2 1 2
2
$D_4$

- *8-Distance ( $\infty$ -Norm)* between  $p_1$  y  $p_2$  :

$$D_8(p_1, p_2) = \max(|x_1 - x_2|, |y_1 - y_2|)$$



2 2 2 2 2
2 1 1 1 2
2 1 0 1 2
2 1 1 1 2
2 2 2 2 2
$D_8$

# Image Convolution

Weighted average of the neighborhood of each pixel

$$g(i, j) = \mathbf{f} \otimes \mathbf{h} = \sum_m \sum_n f(i - m, j - n) h(m, n)$$


The value  $g$  at  $(i, j)$  is the weighted average of  $\mathbf{f}$  at  $(i, j)$  with the values of  $\mathbf{h}$

Equivalent to multiply the coordinates  $(n, m)$  by:


The kernel is flipped in both axes

$$\begin{bmatrix} -1 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} n \\ m \end{bmatrix} = \begin{bmatrix} -n \\ -m \end{bmatrix}$$

Example:







$$g(1,3) = f(0,2) + f(0,1) + f(0,0) + f(1,2) + f(1,1) + f(1,0) + f(2,2) + f(2,1) + f(0,0)$$









# Image Convolution



## How it works (visually):




...



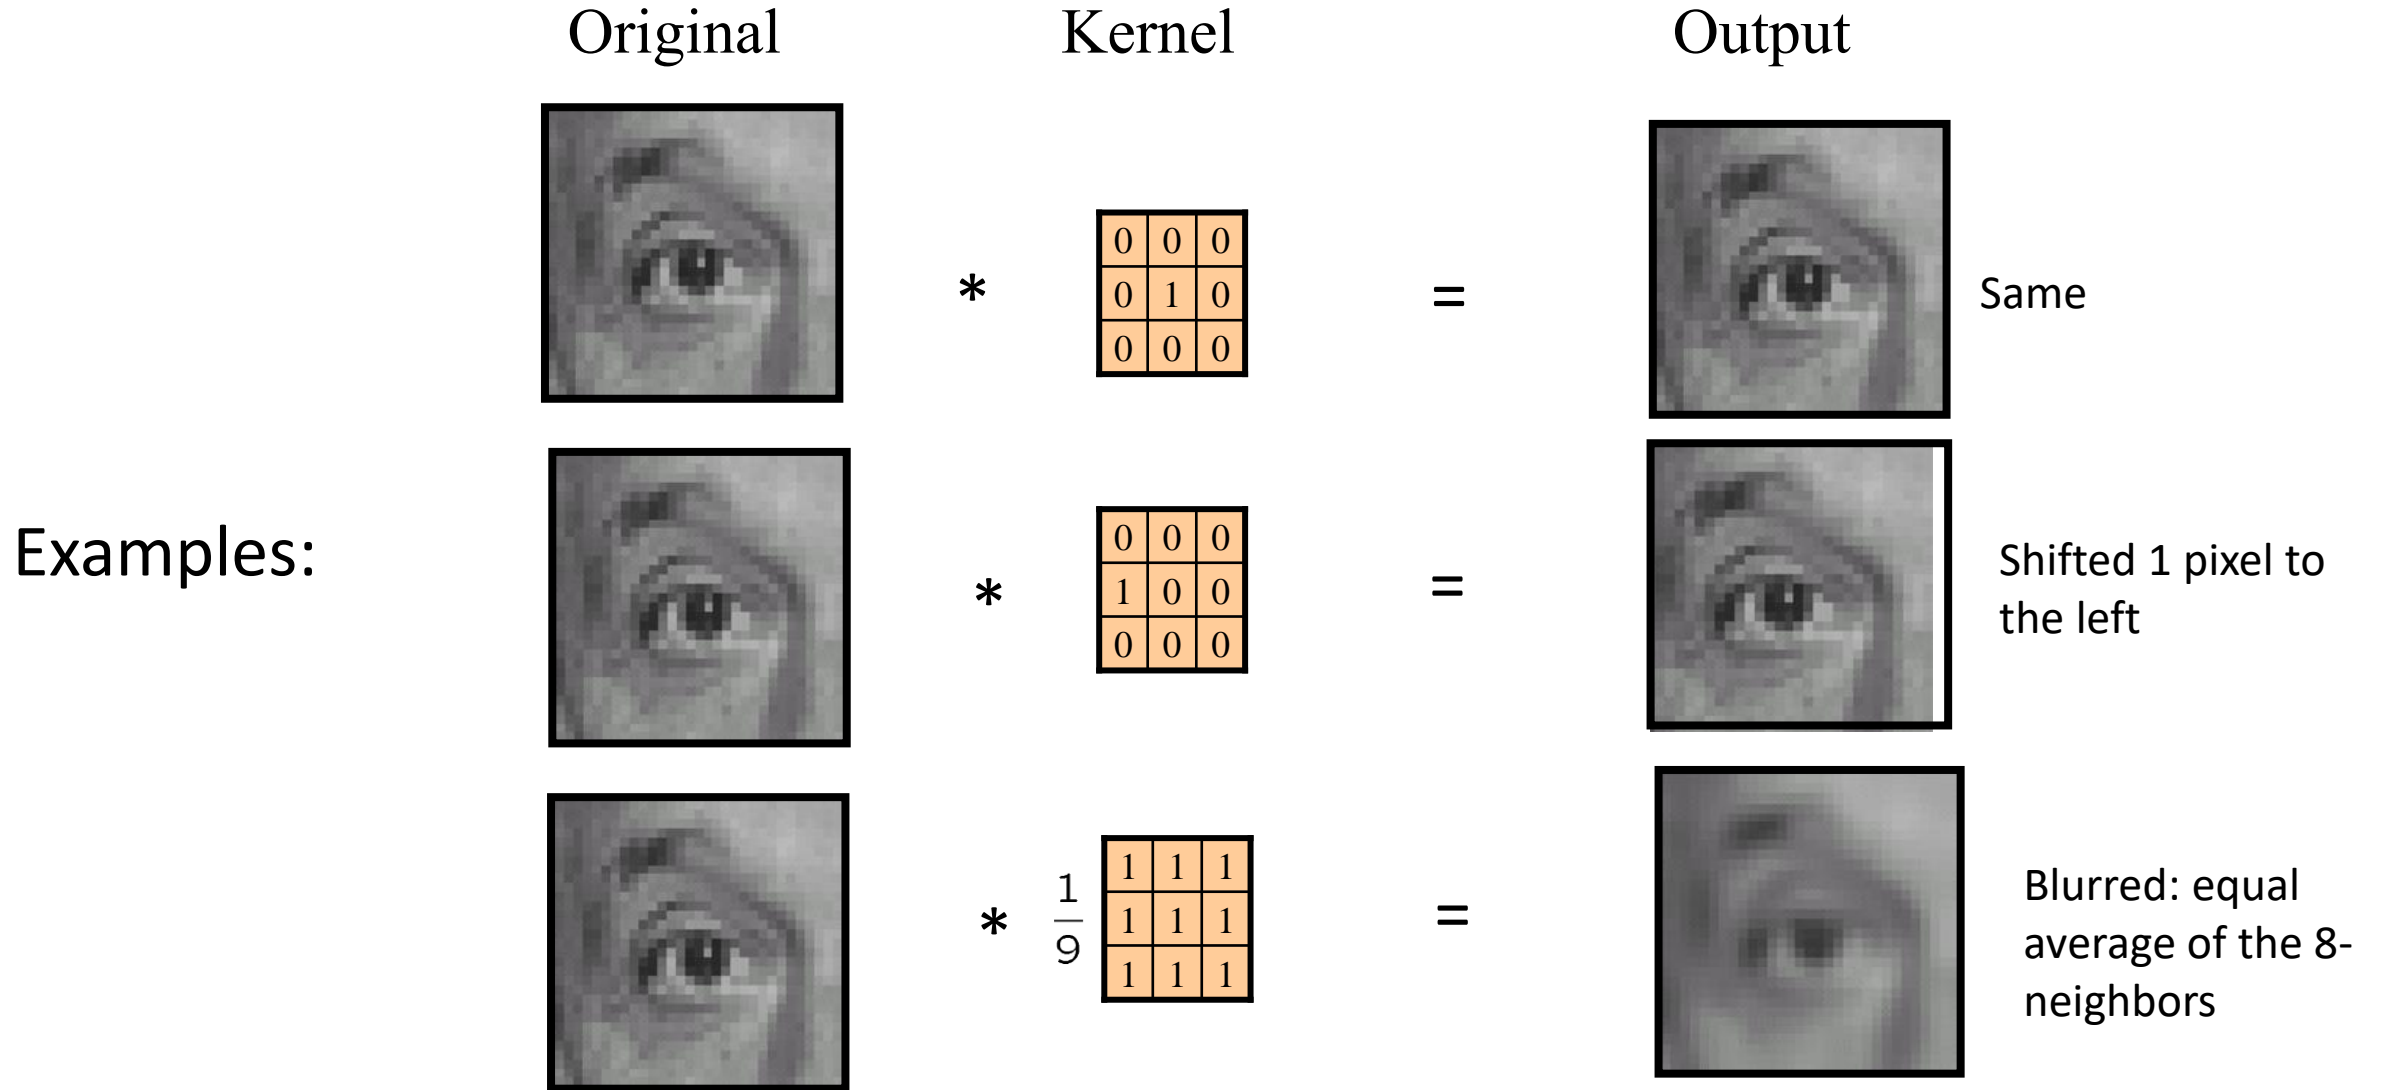
...



\*



# Image Convolution



# Convolution properties

Commutative:  $f * g = g * f$

Associative:  $f * (g * h) = (f * g) * h$

Distributive:  $f * (g + h) = (f * g) + (f * h)$

Asociative wrt scalar product:  $a(f * g) = (af) * g = f * (ag)$

Derivative:  $\mathcal{D}(f * g) = \mathcal{D}f * g = f * \mathcal{D}g$

Theorem of convolution:  $\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$

# 3. Smoothing

Objective: Eliminate/reduce noise in image (in pixel intensities).

Noise appears because of the sensor response (more in CMOS technology), digitalization and transmission.

Approaches:

- Neighborhood averaging
- Gaussian filter
- Median filter
- Image average
- Filters in the frequency domain (not addressed in this course)

# 3. Smoothing

Objective: Eliminate/reduce noise in pixel intensities



Original



Salt and pepper noise



Impulse noise



Gaussian noise

## Salt and pepper noise:

black and white pixels appears at random

## Impulse noise:

black or white pixels appears at random

## Gaussian noise:

- intensities are affected by an additive zero-mean **Gaussian error.**
- simulates thermal noise in sensors, which shows up in **poor illumination conditions**

# 3. Smoothing

## Neighborhood averaging

The value of each pixel is substituted by the average value of the neighbors

$$g(x, y) = \frac{1}{p} \sum_{(m,n) \in S} f(m, n)$$

S : set of "p" pixels in the neighborhood (mxn) of (x,y)

Implemented with the convolution with a kernel (linear operation)

For example, for a 3x3 neighborhood


1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Main **drawback**: edge blur




# 3. Smoothing

## Gaussian noise reduction by Neighborhood averaging



Intensity profile for the row marked in the original image





# Gaussian filter

- is a neighborhood weighted averaging
- weights from a Gaussian bell

$$g_{\sigma}(x,y) = \frac{1}{2\sigma^2\pi} e^{-\frac{1}{2}\left(\frac{x^2+y^2}{\sigma^2}\right)} = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{x}{\sigma}\right)^2} \cdot \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{y}{\sigma}\right)^2}$$

Separability property (rows  
and columns separated)



- Implemented with the convolution of the image with a kernel.

For example 5x5:

0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

# Gaussian filter

Separability property (rows and columns separated)

$$g_{\sigma}(x,y) = \frac{1}{2\sigma^2\pi} e^{-\frac{1}{2}\left(\frac{x^2+y^2}{\sigma^2}\right)} = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{y}{\sigma}\right)^2} \otimes \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{x}{\sigma}\right)^2}$$

$g_{xy}$   $g_y$   $g_x$

Convolution of two 1D kernels

Example:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \otimes \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$


$g_{xy}$   $g_y$   $g_x$

Because of the **Associative property**:

$$\underbrace{f \otimes g}_{\text{2D convolution}} = f \otimes (g_x \otimes g_y) = \underbrace{(f \otimes g_x) \otimes g_y}_{\text{Two 1D convolutions}}$$

# Gaussian filter

$\sigma$  allows us to control the degree of smoothing



## Bigger $\sigma$ ...

- more smoothing
- blurrier image

## Size of the kernel ( $w$ ):

- proportional to  $\sigma$
- must be big enough to account for non-negligible values in the kernel

0	1	3	4	3	1	0
1	6	18	25	18	6	1
3	18	50	71	60	18	3
4	25	71	100	71	25	4
3	18	50	71	60	18	3
1	6	18	25	18	6	1
0	1	3	4	3	1	0

Another row/column would have very small values (negligible)

←  $w$  →

# 3. Smoothing

## Smoothing with **Median filter** (not averaging)

- Replace the intensity value of a pixel by the median of its neighborhood

10	11	12
10	50	10
10	12	11

3x3 image neighborhood around the pixel analyzed (intensity 50)

Sorted list of the neighborhood								
10	10	10	10	11	11	12	12	50

↑  
median

10	11	12
10	11	10
10	12	11

Resulting image

- It is **NOT** a linear operation
- **High computational cost** ... but there are efficient implementations (e.g. pseudomedian, sliding median, ...)

# Median filtering





Image corrupted by  
“salt & pepper” noise



Median




3x3 average

- Preserve borders (**no image blur**)
- Very effective to remove **salt&pepper noise**

# 3. Smoothing

## Image averaging

IDEA: To average several (M) images along a sequence of a static scene (assumption: the only difference is the noise, i.e. static scene)



$$\underbrace{g(x, y)}_{\text{Average image}} = \frac{1}{M} \sum_{i=1}^M f_i(x, y) = \frac{1}{M} \sum_{i=1}^M [f_{\text{noise\_free}}(x, y) + \underbrace{\eta_i(x, y)}_{\text{Noise Image}}] = f_{\text{noise\_free}}(x, y) + \frac{1}{M} \sum_{i=1}^M \eta_i(x, y)$$

If the noise  $\eta_i(x, y)$  has zero-mean  $E[\eta_i(x, y)] = 0$ :      The expected value of noise is zero

$$E[g(x, y)] = E\left[f_{\text{noise\_free}}(x, y) + \frac{1}{M} \sum_{i=1}^M \eta_i(x, y)\right]$$


The expected value is the noise-free image!

$$= f_{\text{noise\_free}}(x, y) + \frac{1}{M} \sum_{i=1}^M E[\eta_i(x, y)] = f_{\text{noise\_free}}(x, y)$$

# 3. Smoothing

## Image averaging

Gaussian noise



1 image



10 images



50 images

Salt&Pepper noise





# 3. Smoothing

Image averaging

Advantage:

- Preserves edges
- Very effective with Gaussian noise (not with salt&pepper)

Drawback:

Only applicable to sequence of images of a still scene

# 4. Image Enhancement

**Objective:** Improve contrast and brightness of the image

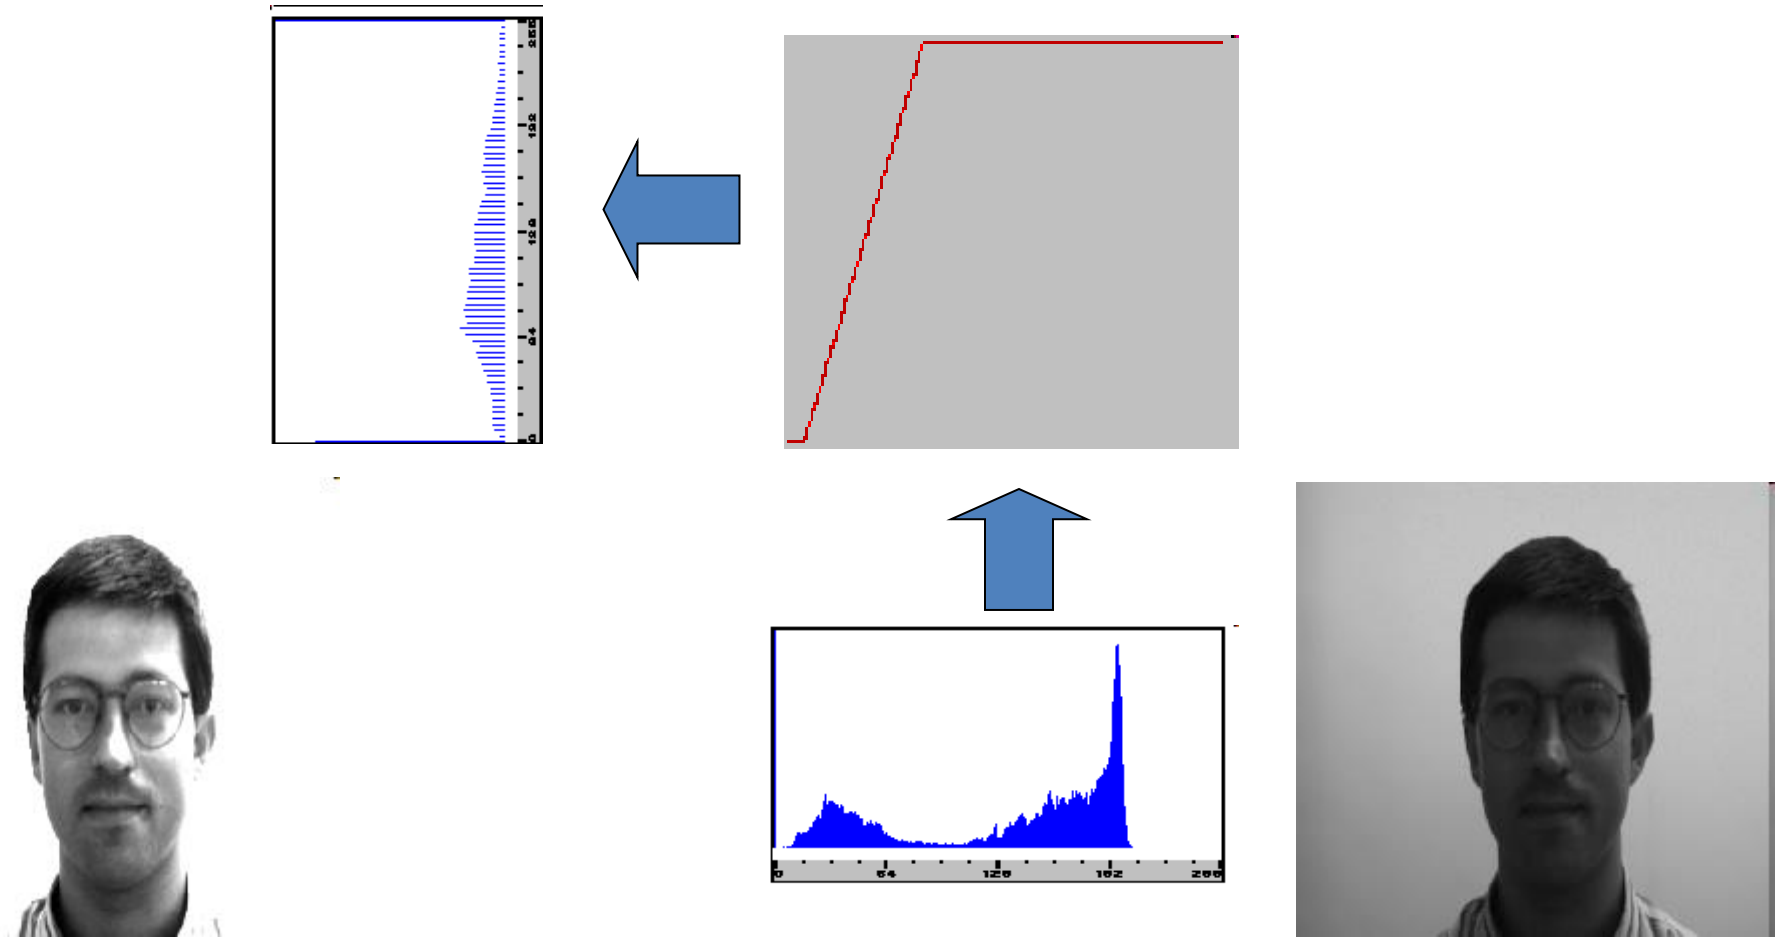
In Computer Vision this is required to **prepare the image for subsequence operations** (feature extraction, segmentation)

Typical operations:

- Lookup-table transformation
- Histogram equalization
- Histogram specification

# Lookup-table transformation


IDEA: To stretch and shift the histogram for a better leverage of the intensity range [typically, 0-255]



# Lookup-table transformation




Non-linear transformation:  $g(x,y) = f(x,y)^\gamma$

- $\gamma < 1$  to increase contrast of darker pixels (Stretching of low intensities)
- $\gamma > 1$  to increase contrast of brighter pixels (Stretching of high intensities)





# Lookup-tables transformation

Non-linear transformation:  $g(x,y) = f(x,y)^\gamma$




$\gamma = 0.5$  (root square)






# Lookup-table transformation

Non-linear transformation:  $g(x,y) = f(x,y)^\gamma$




$\gamma = 3$




# 4. Image Enhancement

## Histogram equalization

IDEA: Modify the intensities such that the new image achieves an uniform histogram




May be applied when low contrast and low brightness (dark) image



# Histogram equalization




Original image



LUT




Equalized image



$h(r_k)$

Histogram of the original image

$$s_k = T(r_k) = \sum_{i=0}^k h(r_i)$$



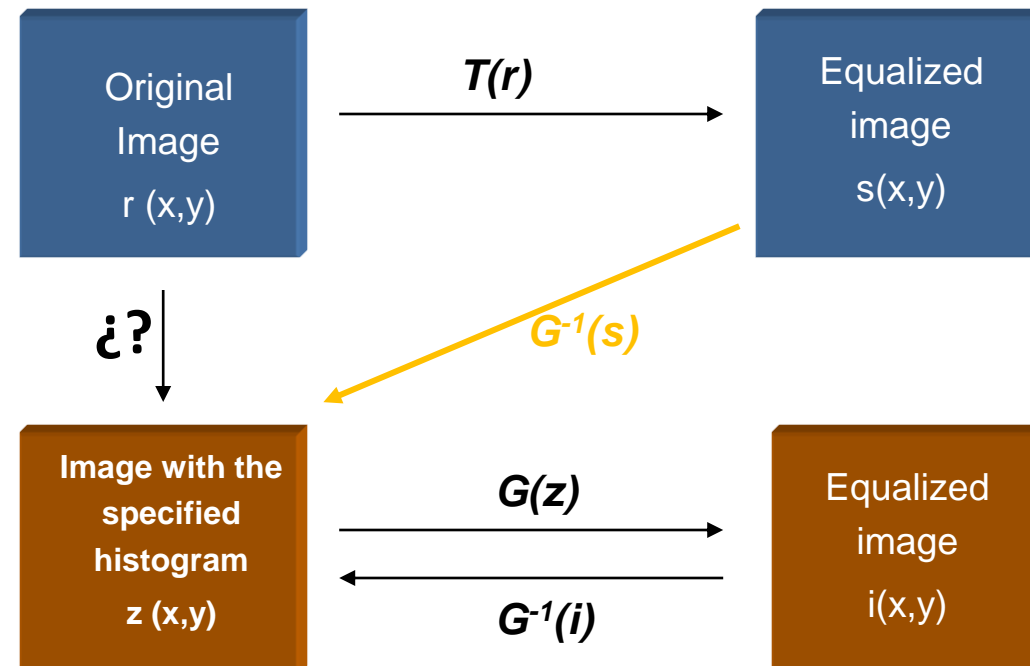
Histogram of the equalized image

In practice, not fully uniform!



# Histogram specification (histogram matching)

- transformation of an image so that its histogram matches a specified histogram
- histogram equalization is a special case in which the specified histogram is uniformly distributed




# Histogram specification

Example:


Original Image




Specified histogram



Enhanced Image



Histogram of the original  
image



Histogram of the enhanced  
image

# Summary

- IP operations needed to improve the quality of the image by attenuating noise, adjusting colors, modifying contrast and brightness, ...
- Necessary to **prepare the image** for a better edge detection or segmentation
- A **color image** is a Tri-dimensional array:  $I(x,y,c)$  with  $c \in \{R, G, B\}$
- A **histogram** provides statistical information of the intensity distribution (e.g. brightness and contrast)
- **LUT transformations** very efficient and effective to change contrast and brightness
- **Convolution** is a linear operation between two images, very useful for IP
- **Gaussian Noise** in images is attenuated by convolving with a Gaussian kernel
- **Salt&Pepper Noise** is attenuated by a median (non-linear) filter