

Edge detection

Javier González Jiménez

Reference Books:

- *Computer Vision: Algorithms and Applications*. Richard Szeliski. Springer. 2010.
<http://szeliski.org/Book>

Content

1. Introduction
2. Operators based on first derivative
3. Operators based on second derivative
4. Canny algorithm

1. Introducción

Edges provide very valuable information about the object contours: useful for segmentation, object recognition, stereo, etc.

750x531x8bit x3
= 1.2Mbytes



1. Introduction

Most of the info of an image is embedded in its edges, and only requires a small proportion of the data




750x531x1bit
= 70 kbytes

1. Introduction

Edges: transitions between two image regions that have very different gray levels (intensities)

Unidimensional, continuous model of an ideal edge:




Parameters

- intensity increment $H = A - B$.
- angle of the slope: “ α ”.
- coordinate “ x_0 ” of the midpoint

1. Introduction

In real images, edges do not follow this model exactly since:



- image are discrete
- are corrupted by noise

The nature of edges may be diverse:

- occlusion borders
- different orientation of surfaces
- different reflectance properties
- different texture
- illumination effects: shadows, reflections, etc.



1. Introduction

Three types of errors involved:

- **Detection error.**

A good detector exhibits low ratio of false negative and false positive.


- **False negatives:** Existing edges are not detected (go unnoticed for the operator)
- **False positives:** Some of the edges detected are no real (induced by noise)

- **Localization error.** Good localizing when the response is at the real, exact position

- **Multiple Response.** Several responses for the same edge




True
edge



Poor robustness
to noise



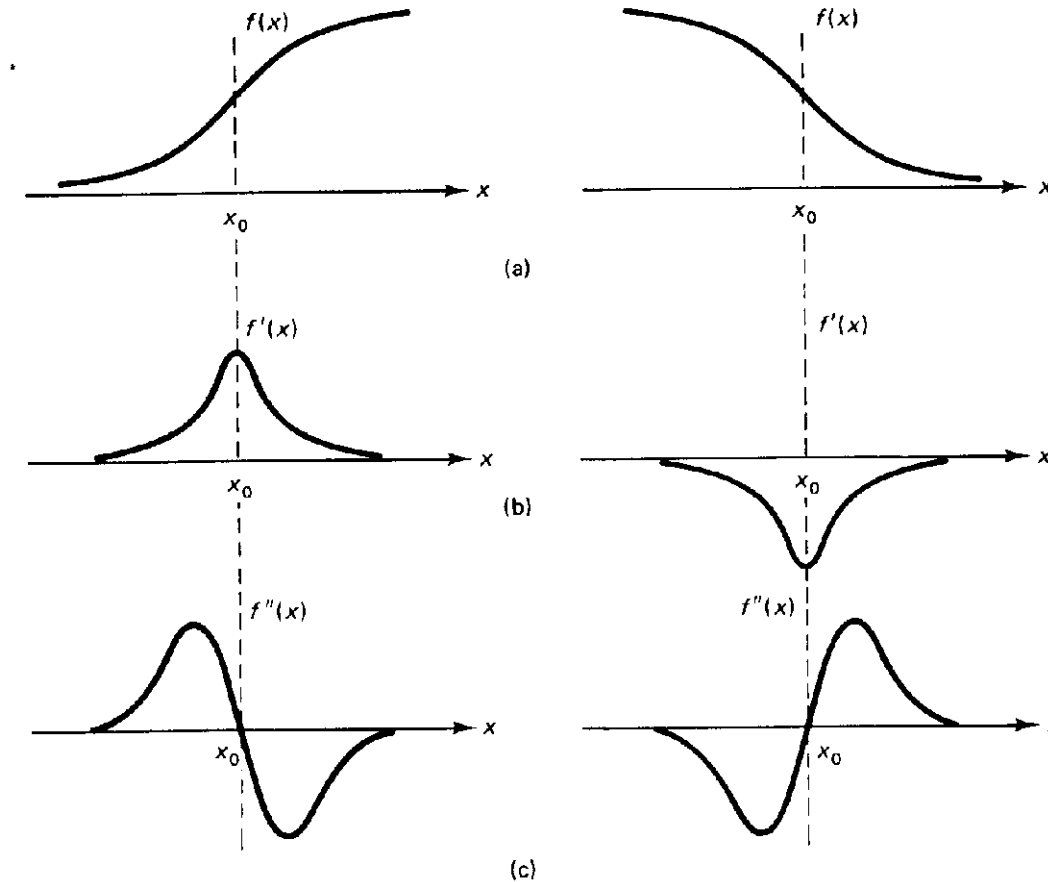
Poor
localization



Too many
responses

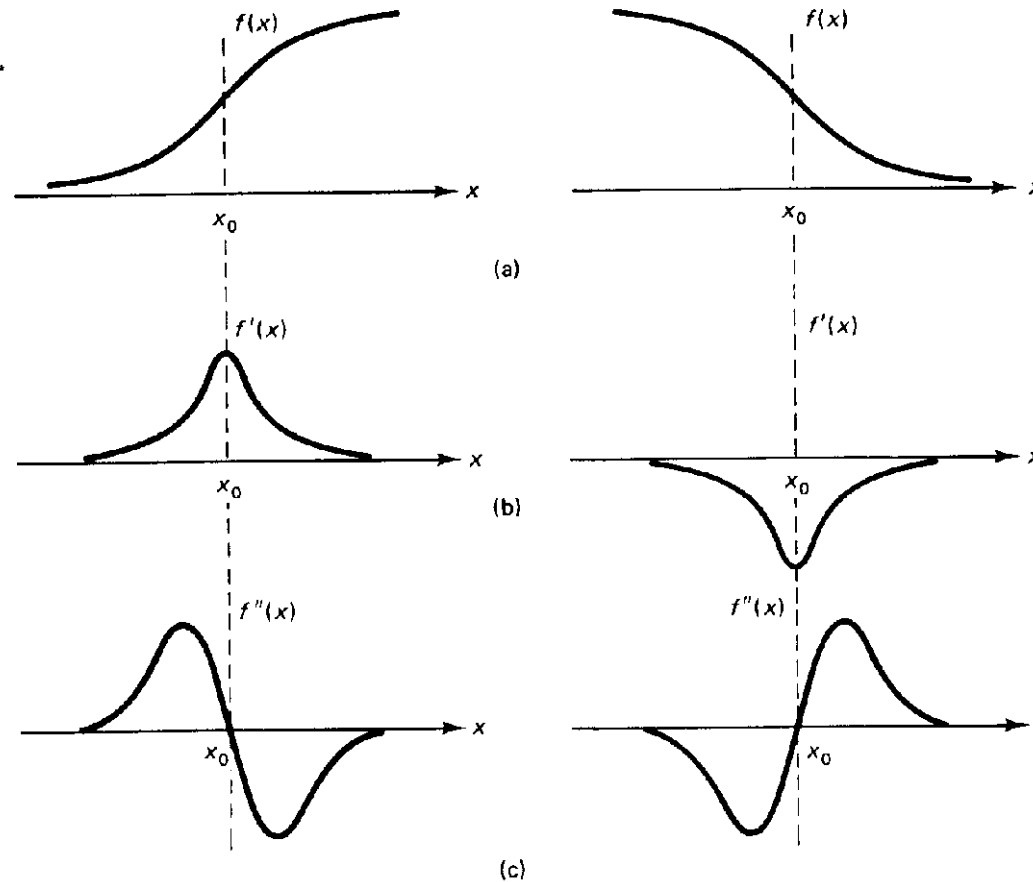
2. Operators based on first derivative (Gradient)

For a one-dimensional continuous function $f(x)$



2. Operators based on first derivative (Gradient)

Intensity transitions are detected through derivatives: For a 1D continuous function $f(x)$



x_0 at the inflection point of $f(x)$

- x_0 at the maximum of $f'(x)$
- The importance of the change is given by $|f'(x_0)|$

x_0 at zero-crossing of $f''(x)$

Fastest variation at x_0

For a **two-dimensional** continuous function $f(x,y)$

Its derivative is a **vector (gradient)**, defined as:


$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial}{\partial x} f(x, y) \\ \frac{\partial}{\partial y} f(x, y) \end{bmatrix} = \begin{bmatrix} f_x(x, y) \\ f_y(x, y) \end{bmatrix}$$

- points at the **direction** of maximum (positive) variation of $f(x,y)$:

$$\alpha(x, y) = \text{atan}\left(\frac{f_y(x, y)}{f_x(x, y)}\right)$$

- **module** proportional to the strength of this variation :

$$|\nabla f(x, y)| = \sqrt{(f_x(x, y))^2 + (f_y(x, y))^2} \approx |f_x(x, y)| + |f_y(x, y)|$$



Discrete approximations of a gradient operator

Based on the differences between gray levels

For example:

- Backward difference of pixels along a row

$$f_x(x, y) \approx G_R(i, j) = [F(i, j) - F(i - 1, j)]/T$$

0	0	0
0	1	-1
0	0	0

- Symmetric difference of pixels along a row

$$f_x(x, y) \approx G_F(i, j) = [F(i + 1, j) - F(i - 1, j)]/2T$$

0	0	0
1	0	-1
0	0	0

- Implemented through the convolution of the image with templates H_R (row) y H_C (column)


$$G_R(i, j) = F(i, j) \otimes H_R(i, j) \quad G_C(i, j) = F(i, j) \otimes H_C(i, j)$$

Discrete approximations of a gradient operator


Roberts	<table><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>-1</td><td>0</td></tr></table>	0	0	0	0	0	1	0	-1	0	<table><tr><td>-1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>0</td></tr></table>	-1	0	0	0	1	0	0	0	0																				
0	0	0																																						
0	0	1																																						
0	-1	0																																						
-1	0	0																																						
0	1	0																																						
0	0	0																																						
Prewitt	$\frac{1}{3}$ <table><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>1</td><td>0</td><td>-1</td></tr></table>	1	0	-1	1	0	-1	1	0	-1	$\frac{1}{3}$ <table><tr><td>-1</td><td>-1</td><td>-1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	-1	-1	-1	0	0	0	1	1	1																				
1	0	-1																																						
1	0	-1																																						
1	0	-1																																						
-1	-1	-1																																						
0	0	0																																						
1	1	1																																						
Sobel	$\frac{1}{4}$ <table><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>2</td><td>0</td><td>-2</td></tr><tr><td>1</td><td>0</td><td>-1</td></tr></table>	1	0	-1	2	0	-2	1	0	-1	$\frac{1}{4}$ <table><tr><td>-1</td><td>-2</td><td>-1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>1</td><td>2</td><td>1</td></tr></table>	-1	-2	-1	0	0	0	1	2	1	$\frac{1}{2+K}$ <table><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>K</td><td>0</td><td>-K</td></tr><tr><td>1</td><td>0</td><td>-1</td></tr></table> $H_F(i,j)$	1	0	-1	K	0	-K	1	0	-1	$\frac{1}{2+K}$ <table><tr><td>-1</td><td>-K</td><td>-1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>1</td><td>K</td><td>1</td></tr></table> $H_C(i,j)$	-1	-K	-1	0	0	0	1	K	1
1	0	-1																																						
2	0	-2																																						
1	0	-1																																						
-1	-2	-1																																						
0	0	0																																						
1	2	1																																						
1	0	-1																																						
K	0	-K																																						
1	0	-1																																						
-1	-K	-1																																						
0	0	0																																						
1	K	1																																						
Frei-Chen	$\frac{1}{2+\sqrt{2}}$ <table><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>$\sqrt{2}$</td><td>0</td><td>$-\sqrt{2}$</td></tr><tr><td>1</td><td>0</td><td>-1</td></tr></table>	1	0	-1	$\sqrt{2}$	0	$-\sqrt{2}$	1	0	-1	$\frac{1}{2+\sqrt{2}}$ <table><tr><td>-1</td><td>$-\sqrt{2}$</td><td>-1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>$\sqrt{2}$</td><td>1</td></tr></table>	-1	$-\sqrt{2}$	-1	0	1	0	1	$\sqrt{2}$	1																				
1	0	-1																																						
$\sqrt{2}$	0	$-\sqrt{2}$																																						
1	0	-1																																						
-1	$-\sqrt{2}$	-1																																						
0	1	0																																						
1	$\sqrt{2}$	1																																						

Example

Original image



(Module) Gradient image




Implementation


↓ $F(i,j)$




Implementation



Implementation



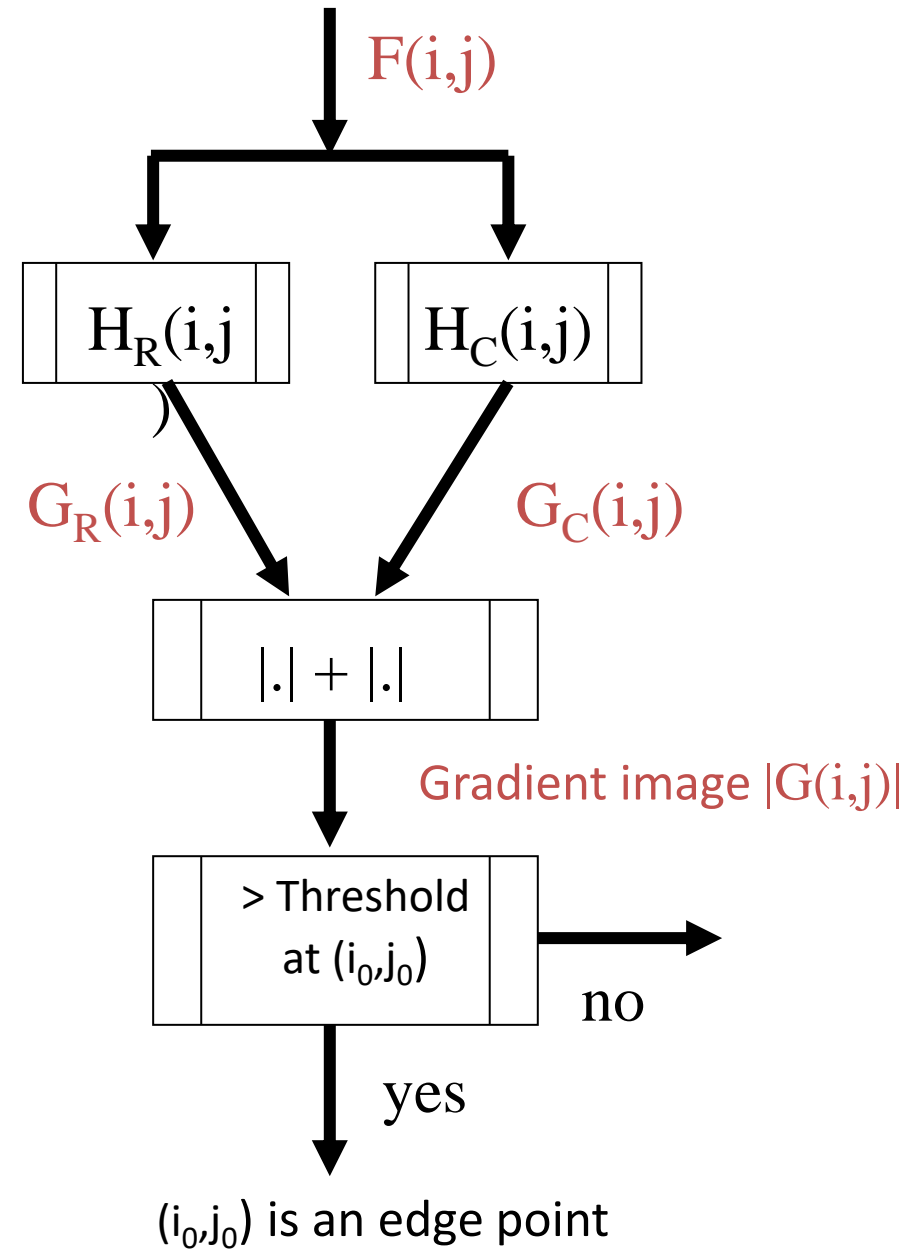
Implementation



Implementation



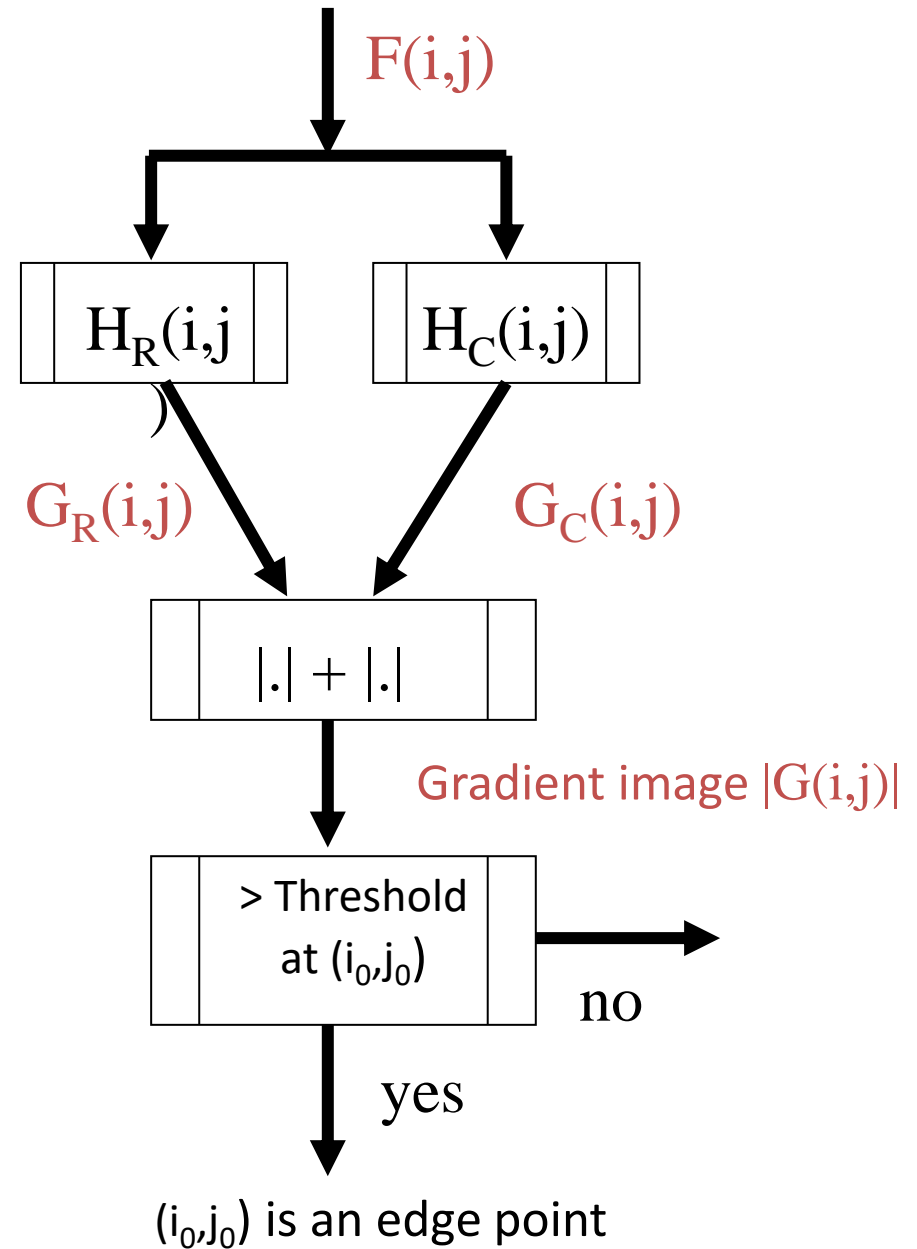
Threshold at 120



Implementation



Threshold at 80



Size of the template ...

Affects the **quality of detection and localization**

Small template →
- more precise localization → good localization
- more affected by noise → likely produces false positive

Large template →
- less precise localization
- more robust to noise → good detector
- higher computational cost $O(N \times N)$


1	0	-1
2	0	-2
1	0	-1

Sobel 3x3

1	2	0	-2	-1
2	3	0	-3	-2
3	5	0	-5	-3
2	3	0	-3	-2
1	2	0	-2	-1

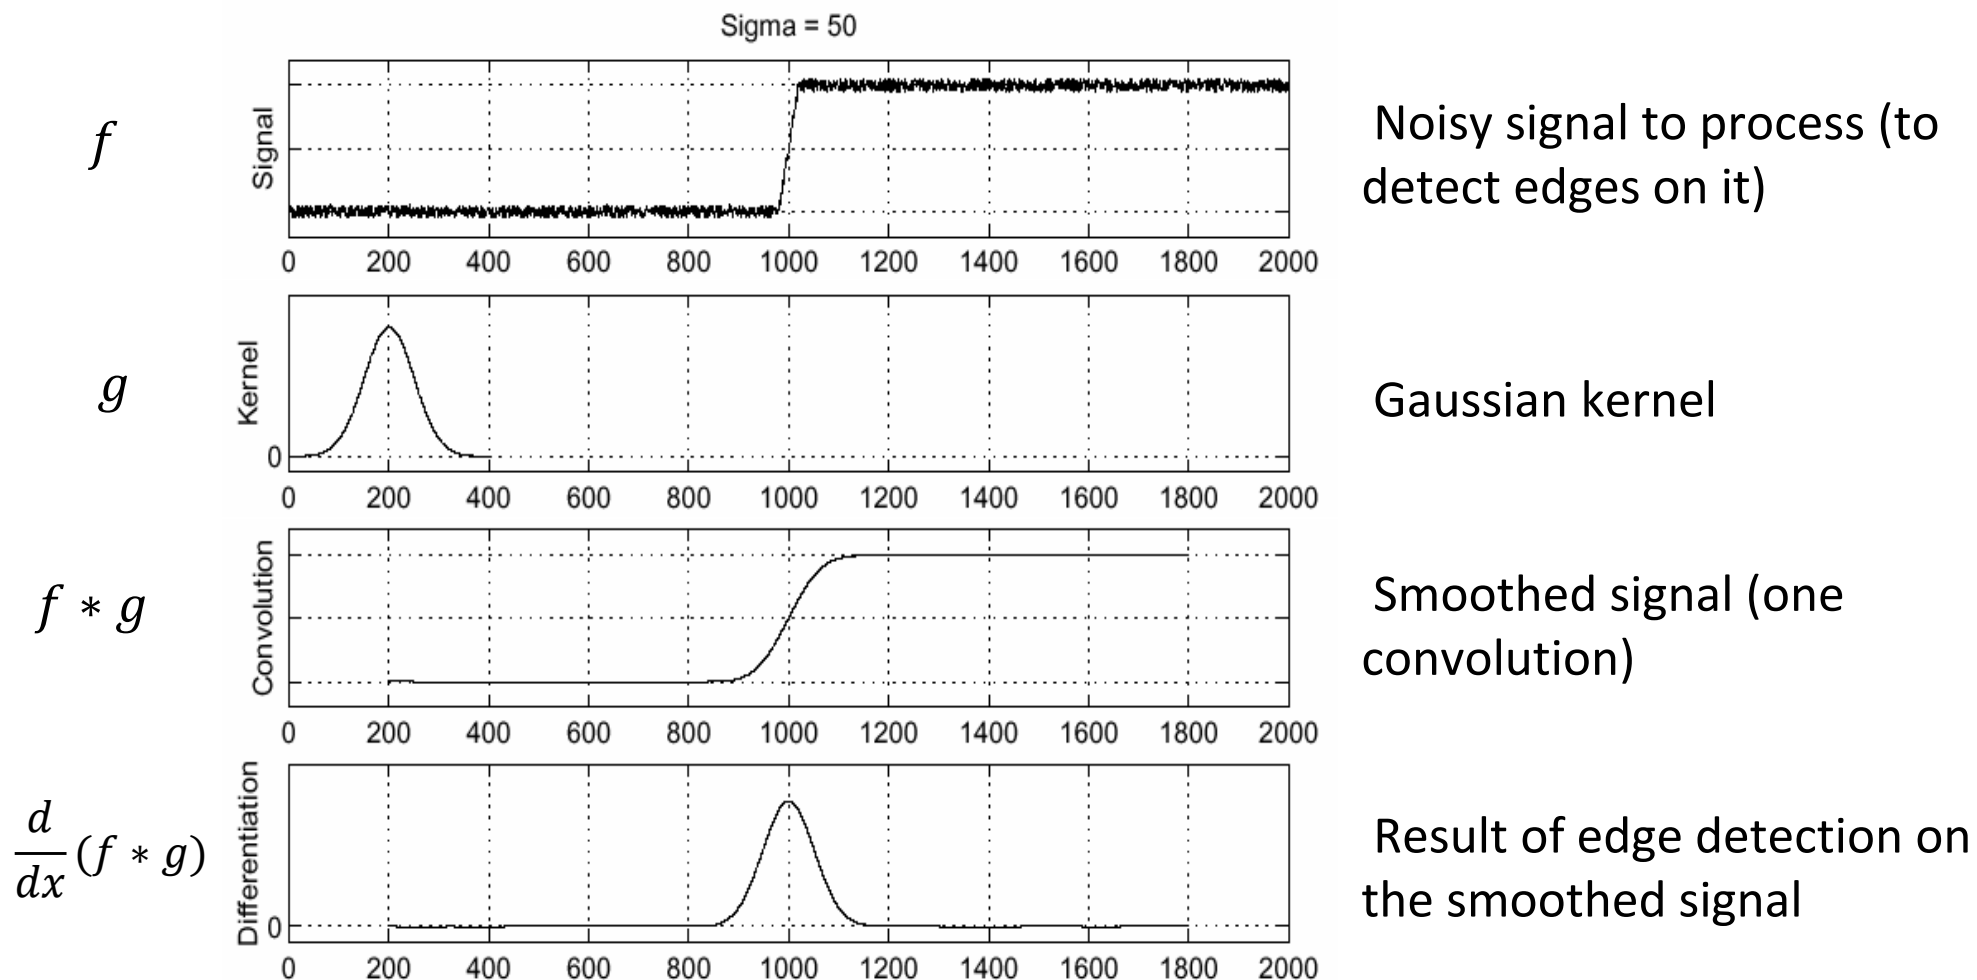
Sobel 5x5

Noise does matter



The response of the derivative to noise is as bigger as the step itself

Solution: smoothing with a Gaussian




To localize edges we have to detect maxima of $\frac{d}{dx}(f * g)$

Smoothing and Gradient combined

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g \quad \text{Convolution property}$$

More efficient if we convolve the image with the derivative of the gaussian (we save one operation)



Smoothing and Gradient combined (2D). **The DroG operator**

DroG (“**D**erivative of **G**aussian”):

- blends smoothing and gradient: $\nabla[f(x,y) \otimes g_\sigma(x,y)]$
- the standard deviation σ controls the degree of smoothness

$$g_\sigma(x,y) = \frac{1}{2\pi\sigma^2} e^{-\frac{1}{2}\frac{(x^2+y^2)}{\sigma^2}} = \frac{1}{2\pi\sigma^2} e^{-\frac{1}{2}\frac{r^2}{\sigma^2}} = g_\sigma(r^2) \quad \begin{array}{l} \text{Radial function.} \\ \text{Not depend on } x,y \end{array}$$

$$\nabla[f(x,y) \otimes g_\sigma(x,y)] = f(x,y) \otimes \nabla[g_\sigma(x,y)] = f(x,y) \otimes \text{DroG}(x,y)$$

separability


$$g(x)' = -xg(x)/\sigma^2$$

$$\text{DroG}(x,y) = \nabla[g_\sigma(x,y)] = \begin{bmatrix} \frac{\partial}{\partial x} [g_\sigma(x)g_\sigma(y)] \\ \frac{\partial}{\partial y} [g_\sigma(x)g_\sigma(y)] \end{bmatrix} = \begin{bmatrix} \frac{-xg_\sigma(x)g_\sigma(y)}{\sigma^2} \\ \frac{-yg_\sigma(x)g_\sigma(y)}{\sigma^2} \end{bmatrix}$$


DroG Operator

Implementation:

- to build a discrete DroG operator, “x, y” are taken in $\mathbb{N} \times \mathbb{N}$
- the DroG template is created just once




DroG along x



DroG along y

$$\text{DroG}(x, y) = \nabla[g_{\sigma}(x, y)] = \begin{bmatrix} \frac{\partial}{\partial x}[g_{\sigma}(x)g_{\sigma}(y)] \\ \frac{\partial}{\partial y}[g_{\sigma}(x)g_{\sigma}(y)] \end{bmatrix} = \begin{bmatrix} \frac{-xg_{\sigma}(x)g_{\sigma}(y)}{\sigma^2} \\ \frac{-yg_{\sigma}(x)g_{\sigma}(y)}{\sigma^2} \end{bmatrix}$$



```
% Separabilidad del operador DroG
% La mascara de derivada en x es: -x*g(x)*g(y)/sigma^2
tamano=5; %Tamaño de la mascara 2xtamano+1
x=[-tamano:tamano];
g_x = fspecial('gaussian',[1,length(x)],2)
g_y = fspecial('gaussian',[length(x),1],2)
drog_x=g_y*(-x.*g_x) %Propiedad de separabilidad
surf(drog_x)
```

DroG Operator

Original image



DroG result
(small σ)



DroG result
(medium σ)




DroG result
(big σ)



3. Operators based on second derivative (Laplacian)

- A second derivative yields a zero-crossing at points where the gradient presents a maximum
- Depending on the edge orientation, this zero-crossing may go almost unnoticed



Solution:

Combine (second) derivatives in perpendicular directions → Laplacian operator

Laplacian

- linear operator
- invariant to image orientation
- precise edge localization

Definition: $\nabla^2 f(i, j) = \frac{\partial^2}{\partial x^2} f(i, j) + \frac{\partial^2}{\partial y^2} f(i, j)$ It's not a vector, but a scalar!

Laplacian is the **trace of the *Hessian matrix***, which fully characterizes the second derivative of a function.

$$\Delta f(x, y) = \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial f}{\partial x} \frac{\partial f}{\partial y} \\ \frac{\partial f}{\partial y} \frac{\partial f}{\partial x} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix}$$

Hessian matrix

Implementation:

First derivative $\frac{\partial f(x, y)}{\partial x} = f_x(x, y) \approx G_F(i, j) = f(i+1, j) - f(i, j)$

$$\frac{\partial^2 f}{\partial x^2} = f_{xx}(x, y) \approx G_F(i, j) - G_F(i-1, j) = f(i+1, j) - 2f(i, j) + f(i-1, j)$$

$$\frac{\partial^2 f}{\partial y^2} = f_{yy}(x, y) \approx G_C(i, j) - G_C(i, j-1) = f(i, j+1) - 2f(i, j) + f(i, j-1)$$

0	0	0
1	-2	1
0	0	0

0	1	0
0	-2	0
0	1	0

Laplacian

Implemented as a convolution: $L[F(i, j)] = F(i, j) \otimes L(i, j)$

0	0	0
1	-2	1
0	0	0

+

0	1	0
0	-2	0
0	1	0

=

0	1	0
1	-4	1
0	1	0

G H $G+H=L$


Distributive property of Convolution

$$f * (g + h) = (f * g) + (f * h)$$

One Convolution

Two Convolutions

An algorithm for zero-crossing detection needed!



Properties

- zero-crossing produced is near a closed contour
- provides edges of 1-pixel width
- poor edge detection because of its sensitivity to noise

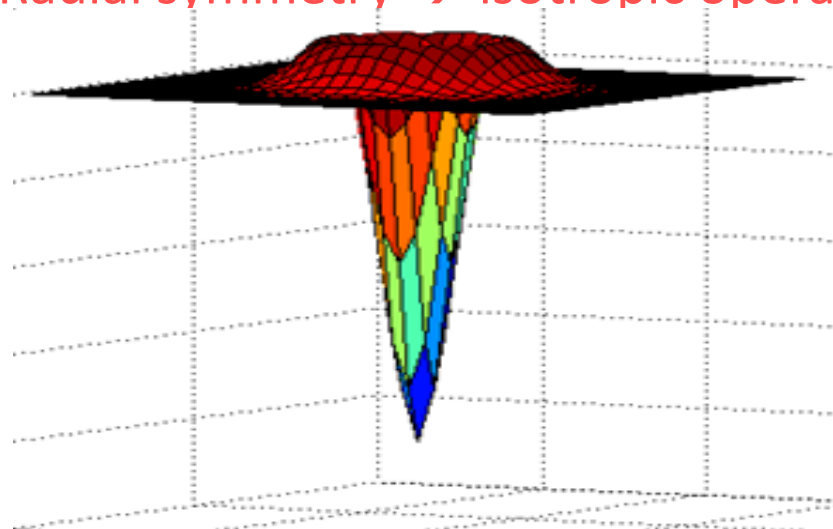
Smoothing and Laplacian combined. LoG operator

- Sensible to noise \rightarrow smoothing needed
- The LoG operator: smoothing and Laplacian (or viceversa, it's commutative)

$$\nabla^2 [f(x, y) \otimes g_\sigma(x, y)] = f(x, y) \otimes \nabla^2 [g_\sigma(x, y)] = f(x, y) \otimes \text{LoG}_\sigma(x, y)$$

$$\text{LoG}_\sigma(x, y) = \frac{1}{\pi\sigma^4} \left[\frac{x^2 + y^2}{2\sigma^2} - 1 \right] e^{-\frac{x^2 + y^2}{2\sigma^2}} = \frac{1}{\pi\sigma^4} \left[\frac{r^2}{2\sigma^2} - 1 \right] e^{-\frac{r^2}{2\sigma^2}} = \text{LoG}_\sigma(r^2)$$


Radial symmetry \rightarrow isotropic operator




applies equally well in all directions in an image

LoG Operator

- High-pass filter (Laplacian) + Low-pass filter (Gaussian) = **Band-pass filter** (width controlled by σ)



LoG in the spatial domain



LoG in the frequency domain

- LoG is **not separable** but can be implemented as **DoG** (“Difference of Gaussians”): $O(N^2) \rightarrow O(4N)$

$$\text{DoG}_{\sigma_1\sigma_2}(x, y) = g_{\sigma_1}(x, y) - g_{\sigma_2}(x, y) = g_{\sigma_1}(x)g_{\sigma_1}(y) - g_{\sigma_2}(x)g_{\sigma_2}(y)$$

Sum of
Separable
Operators

The ratio $\sigma_1/\sigma_2=1.6$ gives the best approximation of LoG

LoG Operator

Drawbacks:

- costly computationally
- not edge orientation is provide
- the output contains negative and non-integer values, so for display purposes the image is normalized to the range 0 – 255
- zero-crossing needed
- tends to round object corners (depending on σ)



LoG Operator

Implementation:

- the image is **convolved with two Gaussians (of different σ)** and the subtracted to each other ($O(2N+2N)$, N being the width of the kernel)
- followed by a **zero crossing algorithm** to detect edges


A simple zero-crossing algorithm

Select a small positive number (threshold) t

A pixel is labelled as edge if in the LoG image


- its value is smaller than $-t$ and at least one of its neighbours is bigger than t ,
or
- Its value is bigger than t and at least one of its neighbours is smaller than $-t$.

LoG Operator



MATLAB


```
I = imread('rice.tif');  
%BW=edge(I,'log',thresh,sigma)  
log = edge(I,'log',0.001,2);  
figure, imshow(log)
```



Open contours

$\sigma=2$
th=0.007


$\sigma=2$
th=0.001



It's an interesting method for **region segmentation** since closed contours provide meaningful regions (of similar pixel intensities)

4. Canny algorithm

- Based on the edge **optimal operator** for a step 1D signal affected with Gaussian noise
- Optimal meeting 3 criteria
 - => good detector
 - => good localizator
 - => single response



CANNY IMPULSE RESPONSE FUNCTIONS

The result is a filter very similar to the DroG in 1D

DERIVATIVE OF GAUSSIAN IMPULSE RESPONSE FUNCTION

4. Canny algorithm

Procedure for images:

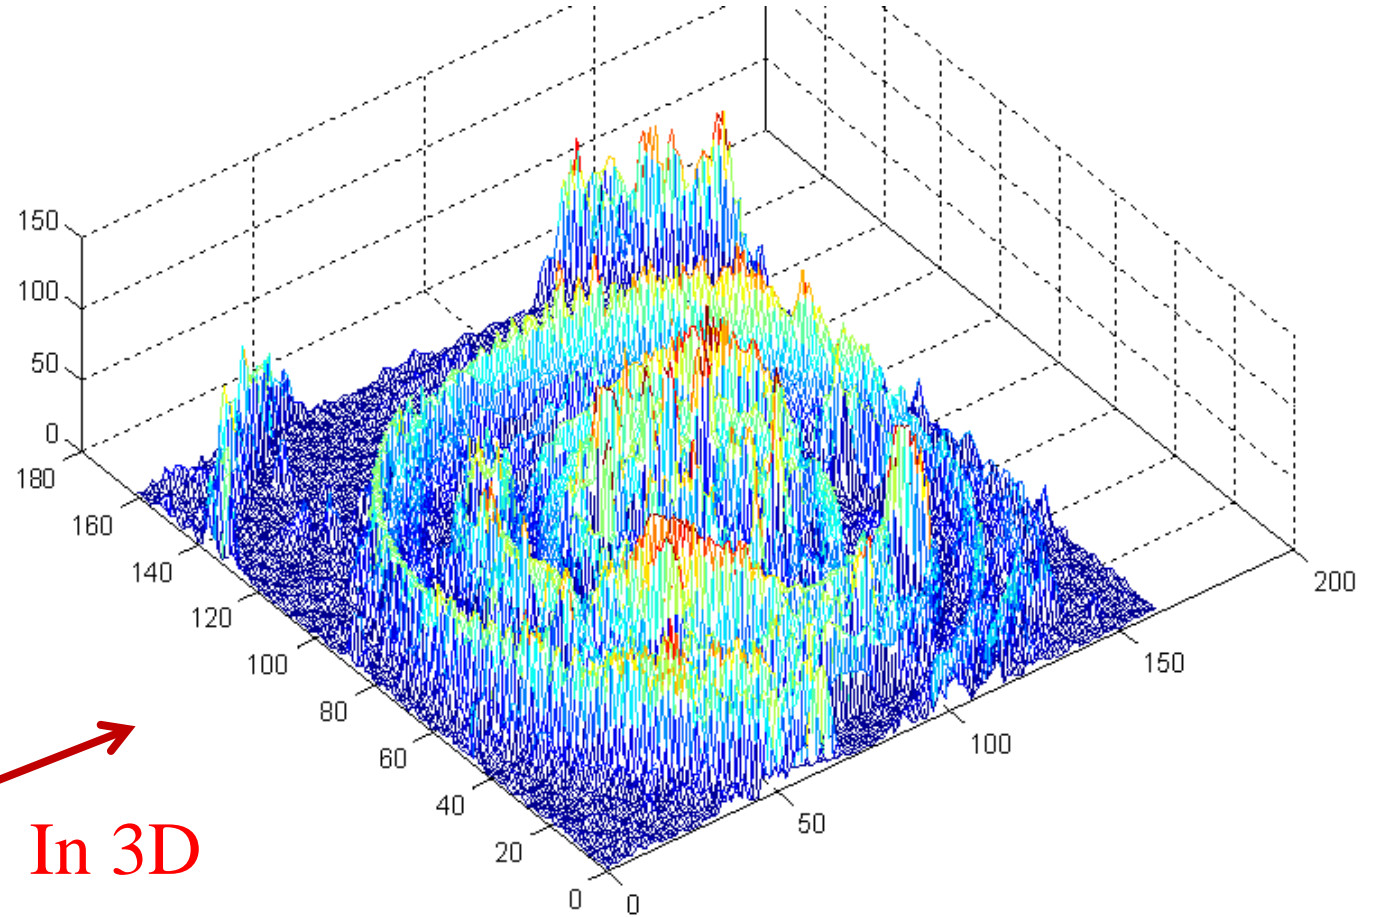

1. Compute the magnitude $G[i,j]$ and the orientation $\theta[i,j]$ of the gradient by applying the **DroG operator**
2. In the direction $\theta[i,j]$ of the gradient vector, apply **nonmaxima suppression**
3. **Threshold** the gradient image with a **hysteresis** to remove less strong edgels (**edge-pixel**) while keeping the contour close

4. Canny algorithm

DroG



DroG





In 3D

We want to follow the summits of these mountains

4. Canny algorithm

Non-maxima suppression

- We consider only the 4 main directions (angular sectors): $[0,45]$, $[45,90]$, $[90,135]$, $[135,180]$.
- The gradient angle $\theta[i,j]$ is approximated to the sector where it lays
- We check $G[i,j]$ at the three points along the selected direction (in red below) and pick the one where $G[i,j]$ is maximum.



We pick the pixel with the biggest where $G[i,j]$


We get the maximum and one single response at each edge

4. Canny algorithm

Why hysteresis: We want to keep the strongest edgels but prevent the disconnection of the contour

With a **single threshold**:


- **Large threshold** → few strong edgels, but disconnected contour
- **Small threshold** → too many edgels (including weak ones)



4. Canny algorithm

Hysteresis


- A high threshold to incorporate strong edges to begin a contour
- A lower threshold to cut-off a contour




4. Canny algorithm

- The algorithm can be repeated with different level of smoothing (changing the sigma of the DroG operator).
- Different sigma produces edges at different spatial scales


Original Image



$\sigma=2$



$\sigma=1$



```
I=imread('mano.bmp');  
imshow(I)  
canny_s1=edge(I,'canny',0.3,1); %BW = edge(I,'canny',thresh,sigma)  
figure,imshow(canny_s1)  
canny_s2=edge(I,'canny',0.3,2);  
figure, imshow(canny_s2)
```

Summary

- Edges are **transitions** between image regions that have different intensities
- An **color image** has 3 edge images! (We can combine them)
- Edges are detected through **derivatives** (smoothing required!):
 - 1st derivative: DroG operator
 - 2nd derivative: LoG/DoG + Zero crossing
- **Canny operator:**
 - Computational expensive (yet in real time!)
 - Very good detection/localization + control of the multiple response
- In all these options, the **Standard Deviation (σ) of the Gaussian** plays an important role → Size of the operator (trade-off between localization and detection)