

Keypoint detection and description

Javier González Jiménez

Reference Books:

- *Computer Vision: Algorithms and Applications*. Richard Szeliski. Springer. 2010.
<http://szeliski.org/Book>

Content

1. Introduction
2. Harris detector
 - Idea
 - Formulation
 - Implementation
3. KLT operator
4. Keypoint matching through correlation
5. SIFT operator
 - Scale Space
 - Detector
 - Descriptor

1. Introduction

What are *keypoints* (also, *interest* or *feature points*)?

Distinctive pixels in the image that can likely be projections on 3D entities.



Synthetic image



Real image

1. Introduction


Two problems involved: Detection and description

1. Detection



The operator must provide a reliable and repeatable response

2. Description: To match to its correspondence in other images



Invariant and discriminative description needed

Harris is just a detector (not a descriptor). The keypoint is described with a surrounding image patch


SIFT is a detector and also provides its own descriptor

1. Introduction

Desired **properties:**

- Detector:**
- Accurate (subpixel accuracy, if possible)
 - Detect all the keypoints in the image (few FALSE NEGATIVES)
 - NOT detect irrelevant points (few FALSE POSITIVES)

Keypoints in the image
(Ground truth)



Detected keypoints
(Prediction)

few FALSE NEGATIVES



Wanted: High **RECALL** (to improve the **SENSITIVITY** of the detector)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Among all matches that are actually positive, what percentage is predicted positive

few FALSE POSITIVES



Wanted: High **PRECISION** (to reduce the false alert rate)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Among all matches that were predicted positive, what percentage is predicted positive

Desired **properties:**

Descriptor should have *invariance* to:

- ✓ Illumination (changes in brightness and contrast)



- ✓ View point (scale, rotación, projective distortion)



Detector and descriptor: Computationally efficient and robust to image noise

1. Introduction

Applications

- Panoramic images (image sticking)
- 3D Reconstruction
- Object tracking
- Object recognition
- Image retrieval and indexing in database
- Robot navigation: mapping, localization, obstacle detection
- and many others

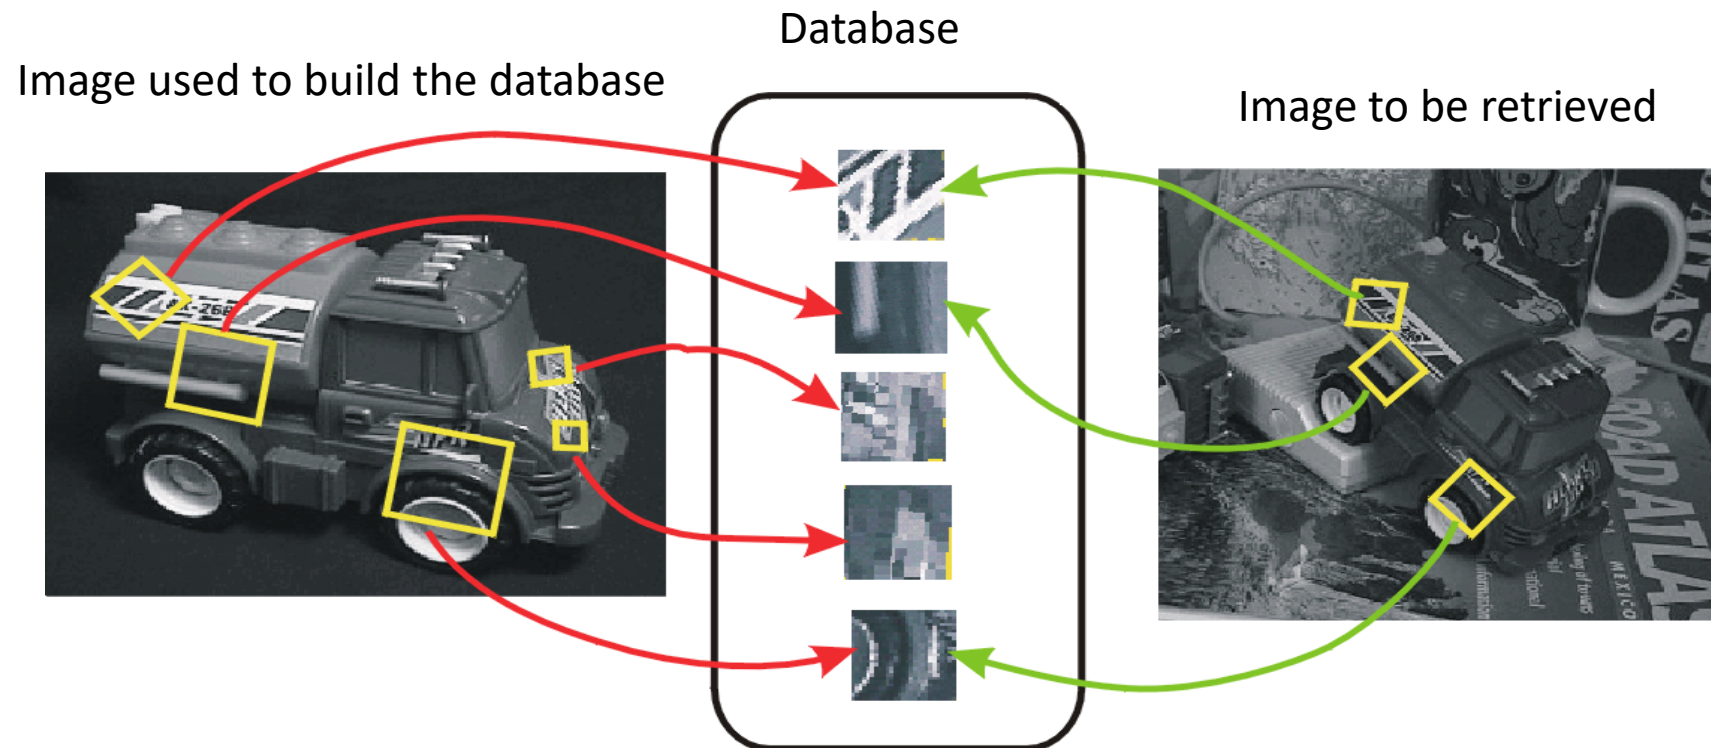
1. Introduction

Applications: Panoramic images (image stitching)




1. Introduction

Applications: Image retrieval and indexing in database




1. Introduction

Applications: Localization and 3D Reconstruction



2. Harris detector


- Detect corners
- Simple and efficient implementation
- Robustness to noise (apply smoothing)
- Invariance to
 - Rotation: uses eigenvectors
 - Brightness (partially to contrast): uses derivatives
- **Not invariance to scale** 



2. Harris detector

- A keypoint is a corner (“esquina”)
- A corner is a point with high variation of intensity in 2 spatial directions

Basic idea: look in a small window W around a pixel if the displacements of the image in two directions provoke changes



Detecting the local change in intensity due to a shift $(\Delta x, \Delta y)$:

Sum-of-square weighted difference at a pixel $[x_0 y_0]$ when a window image I is shifted $(\Delta x, \Delta y)$:

$$E_{x_0 y_0}(\Delta x, \Delta y) = \sum_{x,y} w(x,y) [I(x + \Delta x, y + \Delta y) - I(x,y)]^2$$


Weighting window
centered at $[x_0 y_0]$

Image shifted $(\Delta x, \Delta y)$


Image

RECALL:


1D signal shifting



2D Left-down image shifting




Weighting function $w(x,y)$ may be:



Boolean

1 inside, 0 outside



Gaussian

Gaussian weighting

Let's make more practical the computation of $E_{x_0 y_0}(\Delta x, \Delta y)$

$$E_{x_0 y_0}(\Delta x, \Delta y) = \sum_{x, y} w(x, y) [I(x + \Delta x, y + \Delta y) - I(x, y)]^2 = \sum_{(x_i, y_i) \in W} \underline{[I(x_i + \Delta x, y_i + \Delta y) - I(x_i, y_i)]^2}$$

Sum only over a boolean window W

First order Taylor approximation

Image derivative at (x_i, y_i) along the y axis

$$E(\Delta x, \Delta y) \approx \sum_{(x_i, y_i) \in W} \left[I(\cancel{x_i}, y_i) + [I_x(x_i, y_i) \quad \underbrace{I_y(x_i, y_i)}_{\text{Image derivative at } (x_i, y_i) \text{ along the } y \text{ axis}}] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} - I(\cancel{x_i}, y_i) \right]^2$$

$$= [\Delta x \ \Delta y] \sum_{(x_i, y_i) \in W} \begin{bmatrix} (I_x(x_i, y_i))^2 & I_x(x_i, y_i) I_y(x_i, y_i) \\ I_x(x_i, y_i) I_y(x_i, y_i) & (I_y(x_i, y_i))^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

$$\begin{aligned} (A^T B)^2 &= (A^T B)^T (A^T B) \\ &= B^T A A^T B = B^T M B \end{aligned}$$

$$E_{x_0 y_0}(\Delta x, \Delta y) \approx [\Delta x \ \Delta y] \begin{bmatrix} \sum_W (I_x(x_i, y_i))^2 & \sum_W I_x(x_i, y_i) I_y(x_i, y_i) \\ \sum_W I_x(x_i, y_i) I_y(x_i, y_i) & \sum_W (I_y(x_i, y_i))^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} = [\Delta x \ \Delta y] M \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

M is computed from the **first derivatives at each image point** (x_0, y_0)

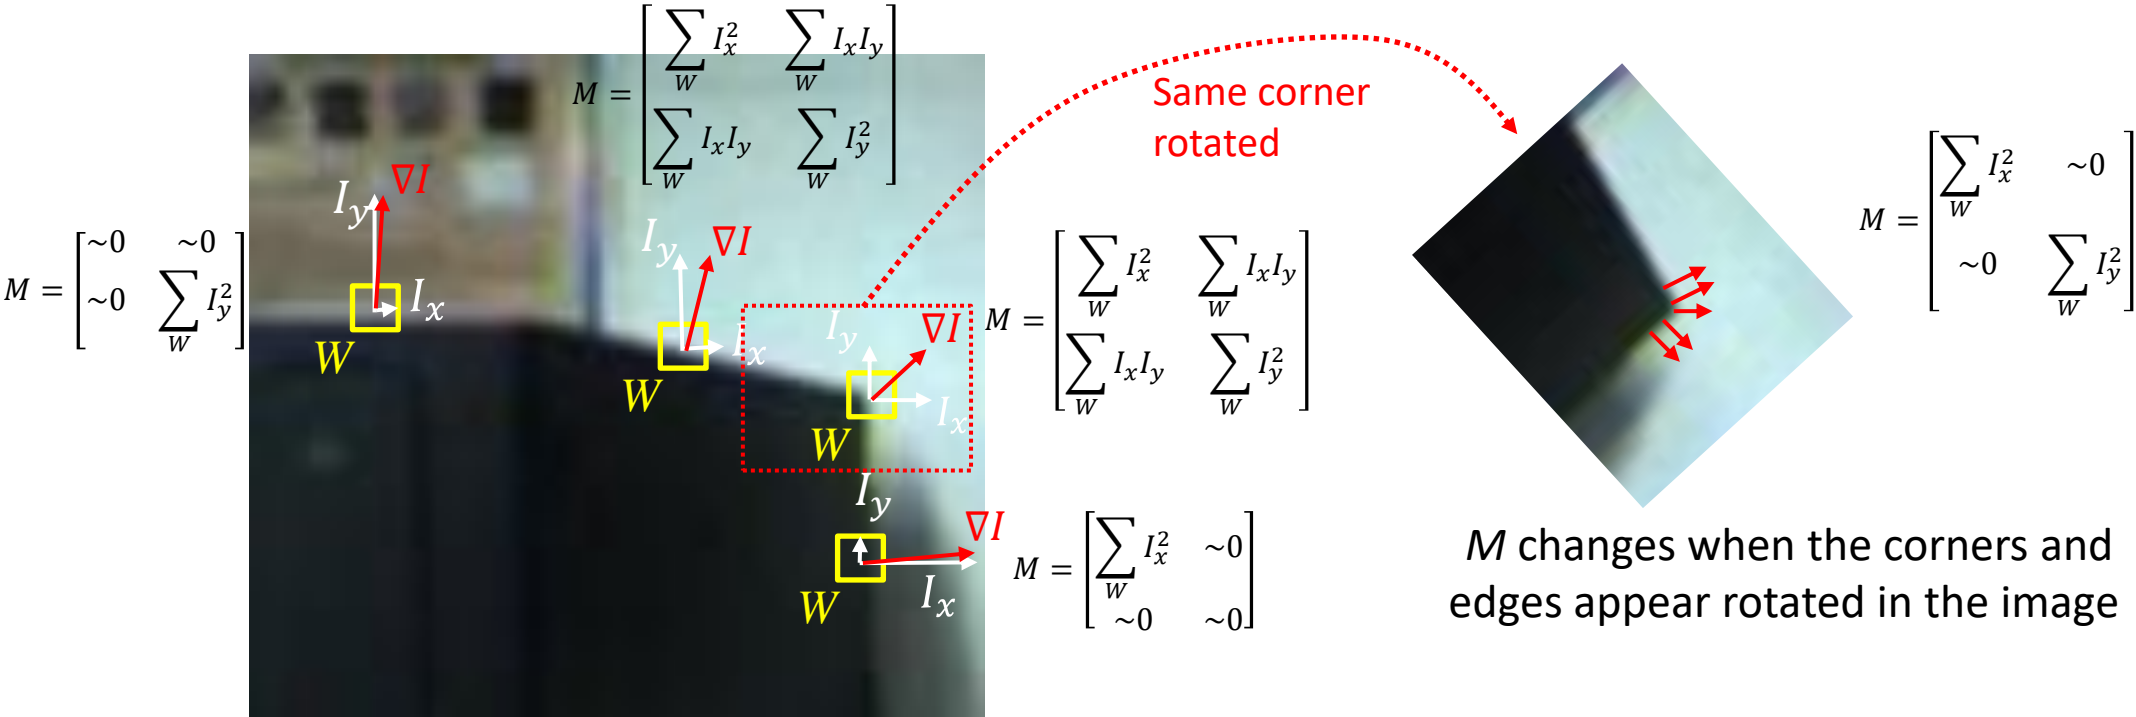
$E_{x_0 y_0}(\Delta x, \Delta y)$ is a **quadratic polynomial** which coefficients are the entries of M

Understanding M

Summation of all the square derivatives
along the x-axis of the points inside W

$$M = \begin{bmatrix} \sum_W I_x^2 & \sum_W I_x I_y \\ \sum_W I_x I_y & \sum_W I_y^2 \end{bmatrix}$$


M is computed from the **first derivatives at each image point (x_o, y_o)**



Derivative vector at different pixels
(always perpendicular to the border)

2. Harris detector

The image points are now classified according to the eigenvalues



But computing the eigenvalues of a 2x2 matrix (M) at each pixel is costly!

2. Harris detector

Let's define a scalar variable R that indexes the same domain:

$$R = \underbrace{\lambda_1 \lambda_2}_{\text{Determinant}} - k \left(\underbrace{\lambda_1 + \lambda_2}_{\text{Trace}} \right)^2 \quad (k = 0.04-0.06 \text{ is an empiric constant})$$

- R is large and positive at corners
- R is negative at edges
- $|R|$ is small at flat regions


Trace and determinant of M and $D = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$ are the same:

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

$$R = \det M - k (\text{trace } M)^2$$





No need to compute the eigenvalues!!



2. Harris detector

Summary:





Original image



2. Harris detector

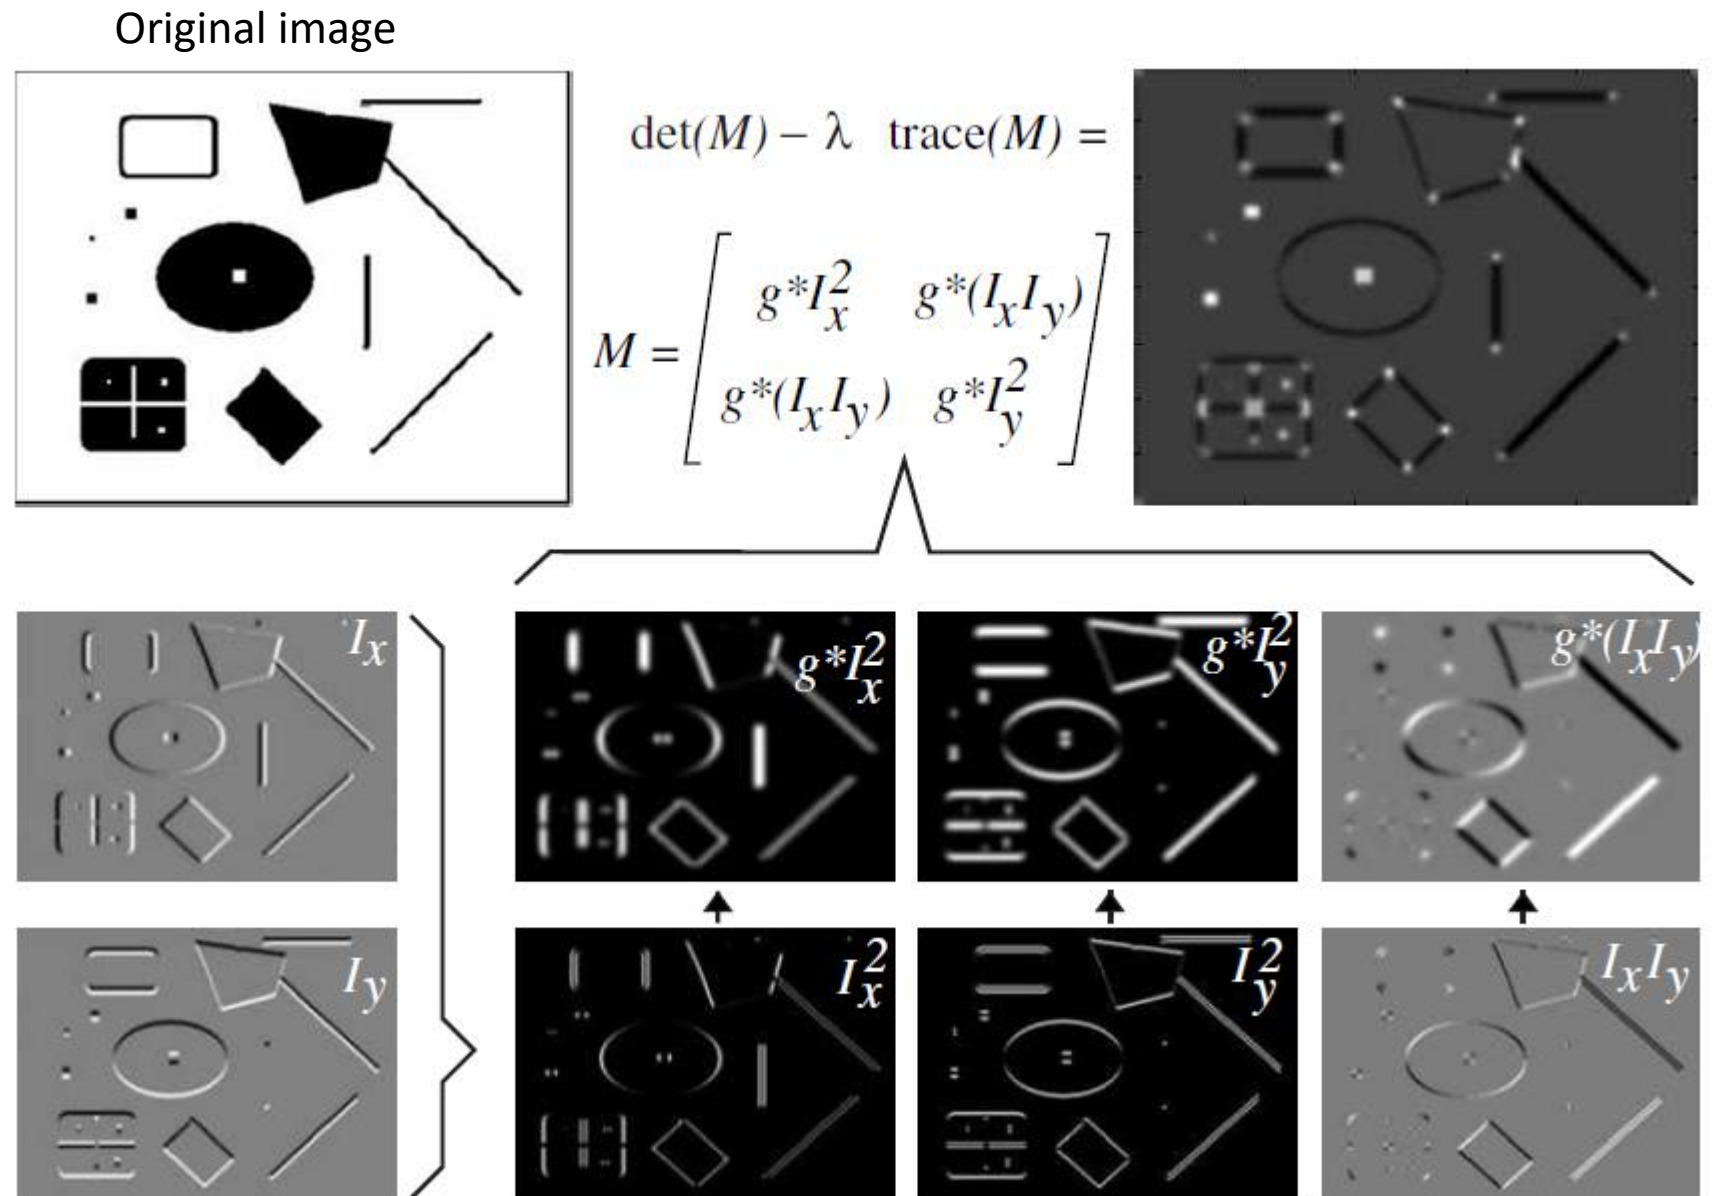
Summary:

Original image



2. Harris detector

Summary:



2. Harris detector

Algorithm:


- Compute the image derivatives I_x , I_y (i.e. Sobel)
- Gaussian smoothing of the 3 images: $(I_x)^2, (I_y)^2, I_x I_y$
- Compute the image R from the formula (trace and determinant)
- Find pixels where R is high ($R > \text{threshold}$)
- Select local maxima

2. Harris detector




2. Harris detector

R image



2. Harris detector

$R > \text{threshold}$



2. Harris detector

Local maxima of R



2. Harris detector



3. KLT operator (Kanade-Lucas-Tomasi)

Objective: Detect distinctive points, suitable to be tracked in a image sequence


- Similar principle to Harris: “A good keypoint is that with a high intensity derivative in two directions” → **Min (λ_1, λ_2) > threshold**
- Also based on the first derivative matrix:

$$M = \begin{bmatrix} \sum_w I_x^2 & \sum_w I_x I_y \\ \sum_w I_x I_y & \sum_w I_y^2 \end{bmatrix} \longrightarrow D = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

- But now, the **eigenvalues are computed** (no approximation with R)
→ better behavior under affine image deformation

4. Keypoint matching through correlation

Which is the correspondence of each point in other image?



Sum of squared differences (SSD): $SSD(f, h) = \sum [f(i, j) - h(i, j)]^2$

SSD is approximated by the SAD (more efficient computationally):

Sum of absolute differences (SAD): $\sum |f(i, j) - h(i, j)|$



Problem: SSD and SAD are not invariant to brightness or contrast changes

Still, SAD is employed in keypoint tracking along an image sequence (where image brightness and contrast do not change very much).

4. Keypoint matching through correlation

Cross-Correlation:


$$CC(x, y) = \sum_{i,j} f(x + i, y + j)h(i, j)$$



CC is similar to the convolution but without flipping the kernel

$$f \otimes h = \sum_{x,y} f(x - i, y - j)h(i, j) \quad \text{equivalent to the cross-correlation of } f(-i, -j) \text{ and } h(i, j)$$

Correlation vs. Convolution:



The convolution $f * g$ is equivalent to the cross-correlation of $f(t)$ and $g(-t)$

Normalized Cross-Correlation (NCC):

Cross correlation is not invariant to change in brightness and contrast of f & h
→ Normalization required

$$NCC(x, y) = \sum_{i,j} \hat{f}(x+i, y+j) \hat{h}(i, j)$$

Normalization: \hat{f} and \hat{h} have **zero mean** and **contrast one**

$$\hat{f}(x+i, y+j) = \frac{f(x+i, y+j) - \bar{f}}{\|f - \bar{f}\|_{W_m(x,y)}} \quad \text{brightness} \quad \hat{h}(x, y) = \frac{h(x, y) - \bar{h}}{\|h - \bar{h}\|_{W_m(x,y)}} \quad \text{brightness}$$

Mean brightness (intensity) of f and h in the window W_m

$$\bar{f} = \frac{1}{|W_m(x, y)|} \sum_{(i,j) \in W_m(x,y)} f(i, j)$$

$$\bar{h} = \frac{1}{|W_m(x, y)|} \sum_{(i,j) \in W_m(x,y)} h(i, j)$$

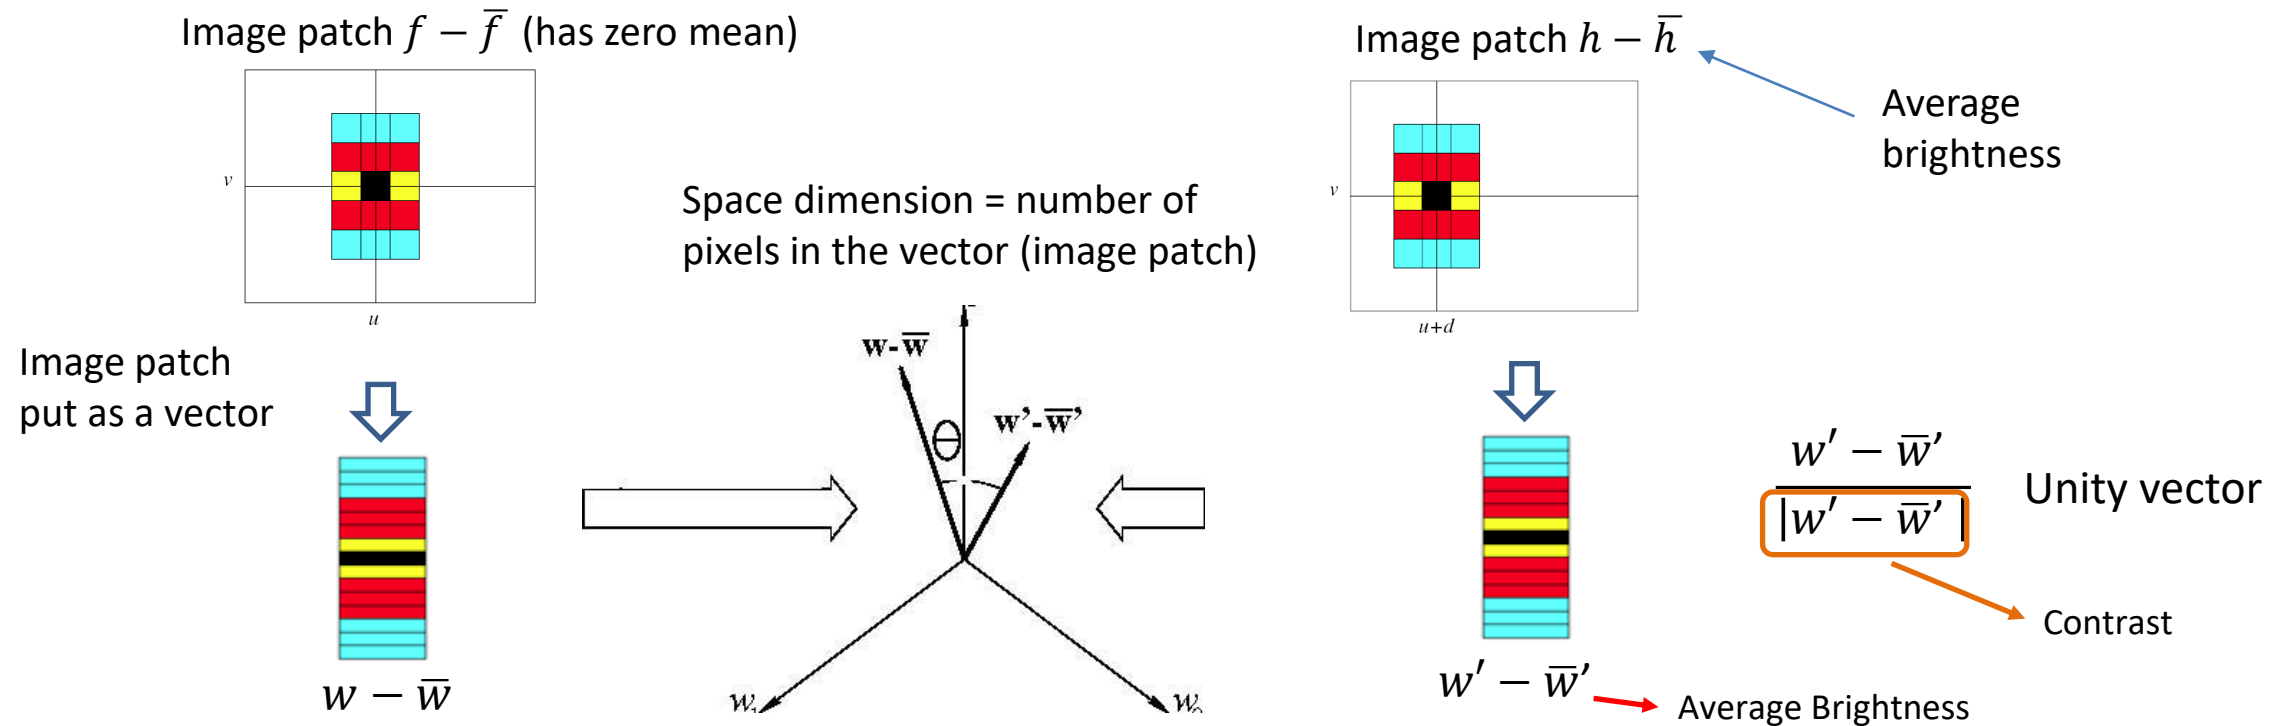
Contrast (norm) of brightness in the window W_m

$$\|f - \bar{f}\|_{W_m(x,y)} = \sqrt{\sum_{(i,j) \in W_m(x,y)} [f(i, j) - \bar{f}(i, j)]^2}$$

$$\|h - \bar{h}\|_{W_m(x,y)} = \sqrt{\sum_{(i,j) \in W_m(x,y)} [h(i, j) - \bar{h}(i, j)]^2}$$

Why does NCC measure similarity between two image patches?

Let's consider an image patch as a vector



$$C(d) = \frac{1}{|w - \bar{w}|} \frac{1}{|w' - \bar{w}'|} (w - \bar{w}) \cdot (w' - \bar{w}') = \cos \theta$$

The similarity between two unity vectors is given by the angle (or \cos) between them


4. Keypoint matching through correlation



*




=



*



=



Correlation output

DEMO IN MATLAB

```
%Read and show the image
flowers = imread('flowers.tif'); figure, imshow(flowers)
% select a template from the image with the mouse
[sub_flowers,rect_flowers] = imcrop(flowers);
% Show the selected template
figure, imshow(sub_flowers)
% Do NCC with the blue channel and display the result
c = normxcorr2(sub_flowers(:,:,1),flowers(:,:,1));
figure, surf(c), shading flat
```


Notice: Output not invariant to the rotation of the patch

Content

1. Introduction
2. Harris detector
 - Idea
 - Formulation
 - Implementation
3. KLT operator
4. Keypoint matching through correlation
5. SIFT operator
 - Scale Space
 - Detector
 - Descriptor

5. The SIFT (**Scale Invariant Feature Transform**) operator

- **Objective:** Find in image projections of **distinctive** 3D points (not necessary corners!) that are **invariant to scale**



- **Provides** both: Detector y descriptor of the detected keypoints

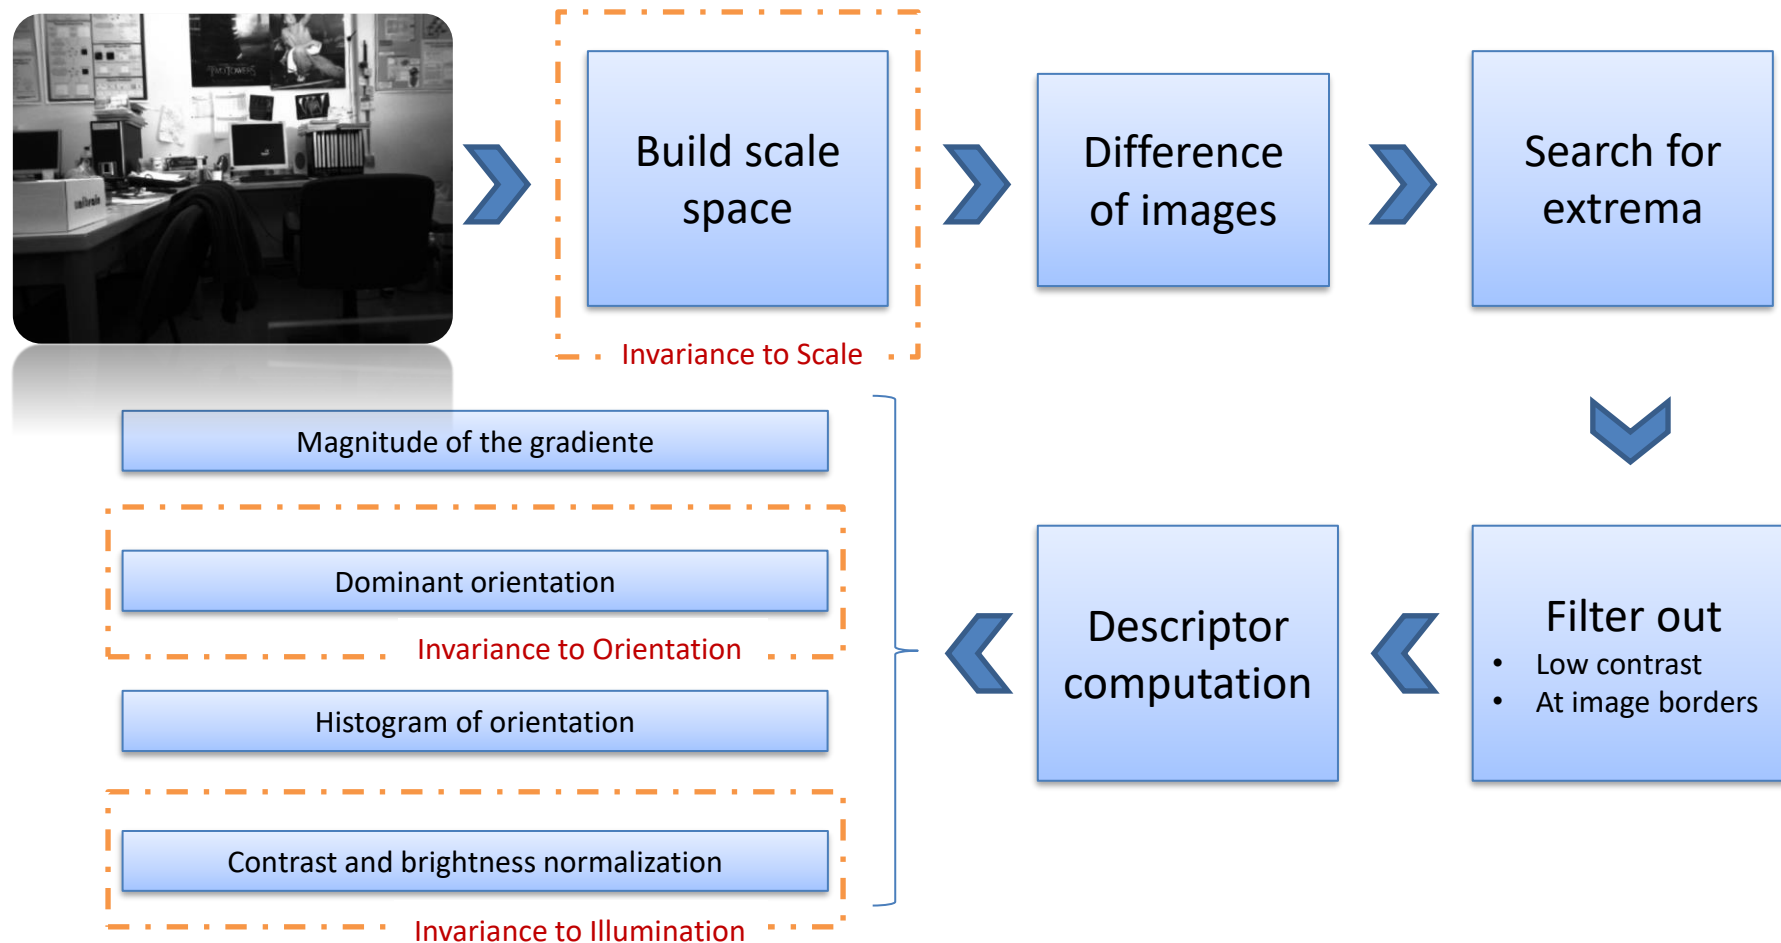
Proposed (and patented) by **David Lowe**: "**Distinctive image features from scale-invariant keypoints**," *International Journal of Computer Vision*, 60, 2 (2004).

5. The SIFT operator

- Both, detector and descriptor have invariance to:
 - Scale Important difference against Harris
 - Rotation Important difference against NCC
 - Illumination [contrast+brightness]
 - Affine transformation [parcially]
- Principle
 - Search for extrema in the scale space [**Detector**]
 - Normalized histogram of orientation [**Descriptor**]
- Descriptor is a vector up to 128 dimensions

5. The SIFT operator


Overview of the method



5. The SIFT operator

Scale Space

In images, features emerge at different scales



Low resolution

Scale

High resolution

Changing the scale (size) shows up different features




5. The SIFT operator

Scale Space

Features show up at a given scale


Example in one-dimension



5. The SIFT operator

Scale Space


We can change the scale by **smoothing the signal with a Gaussian**



5. The SIFT operator

Scale Space

Changing the scale of an image by **smoothing with a Gaussian**



Scale Space

Changing the scale of an image by **smoothing with a Gaussian**

Gaussian operator: $G(x, y, t) = \frac{1}{2\pi t} e^{-\frac{(x^2+y^2)}{2t}} \quad t = \sigma^2 \Rightarrow \sigma = \sqrt{t}$

Smoothed image: $L(x, y, t) = I(x, y) * G(x, y, t)$



$$L(x, y, 0) = I(x, y)$$

starting image $t=0$



$$L(x, y, 1)$$

$\sigma=1$



$$L(x, y, 4)$$

$\sigma=2$

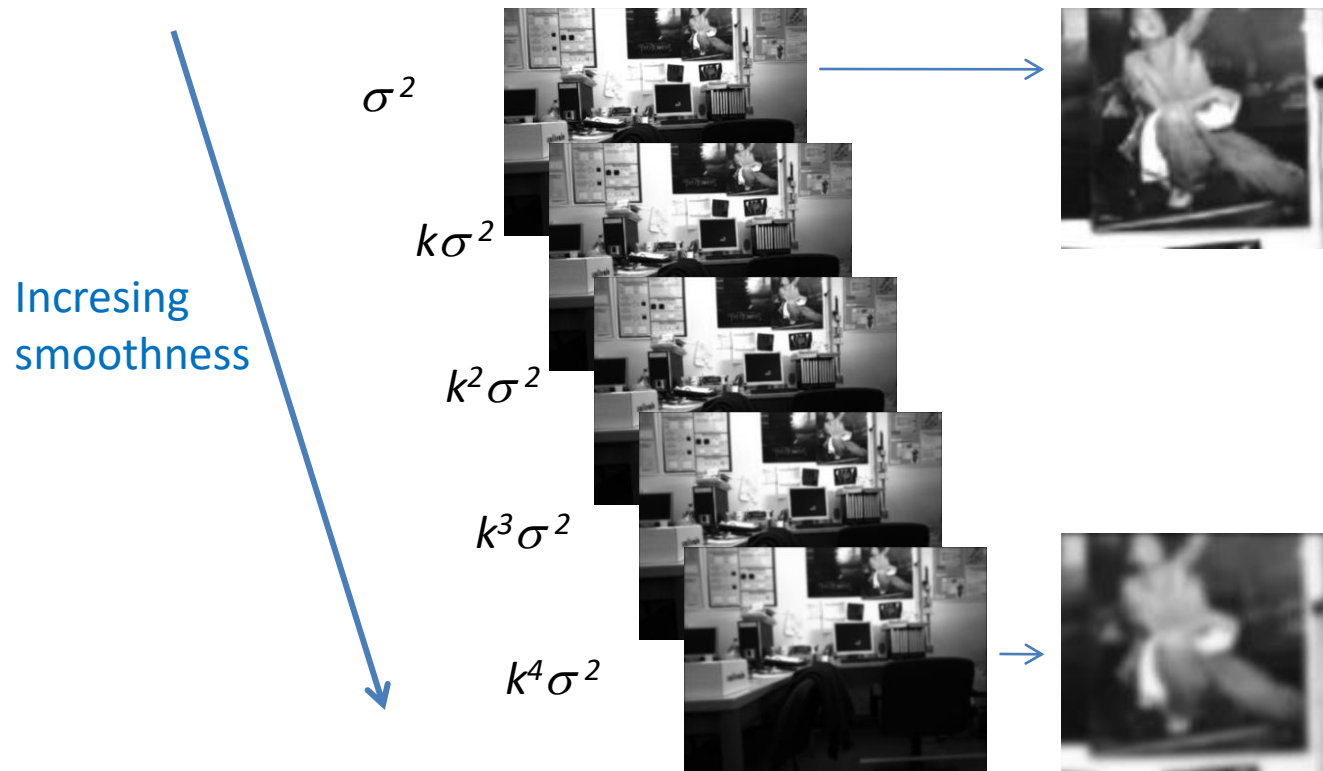
The starting image is always the original image ($t=0$)

So, the scale space stores samples of the function $L(x, y, \sigma^2)$

5. The SIFT operator


Objective: “continuous” scale space

Progressive convolution of the input image with a Gaussian controlled with a constant factor K



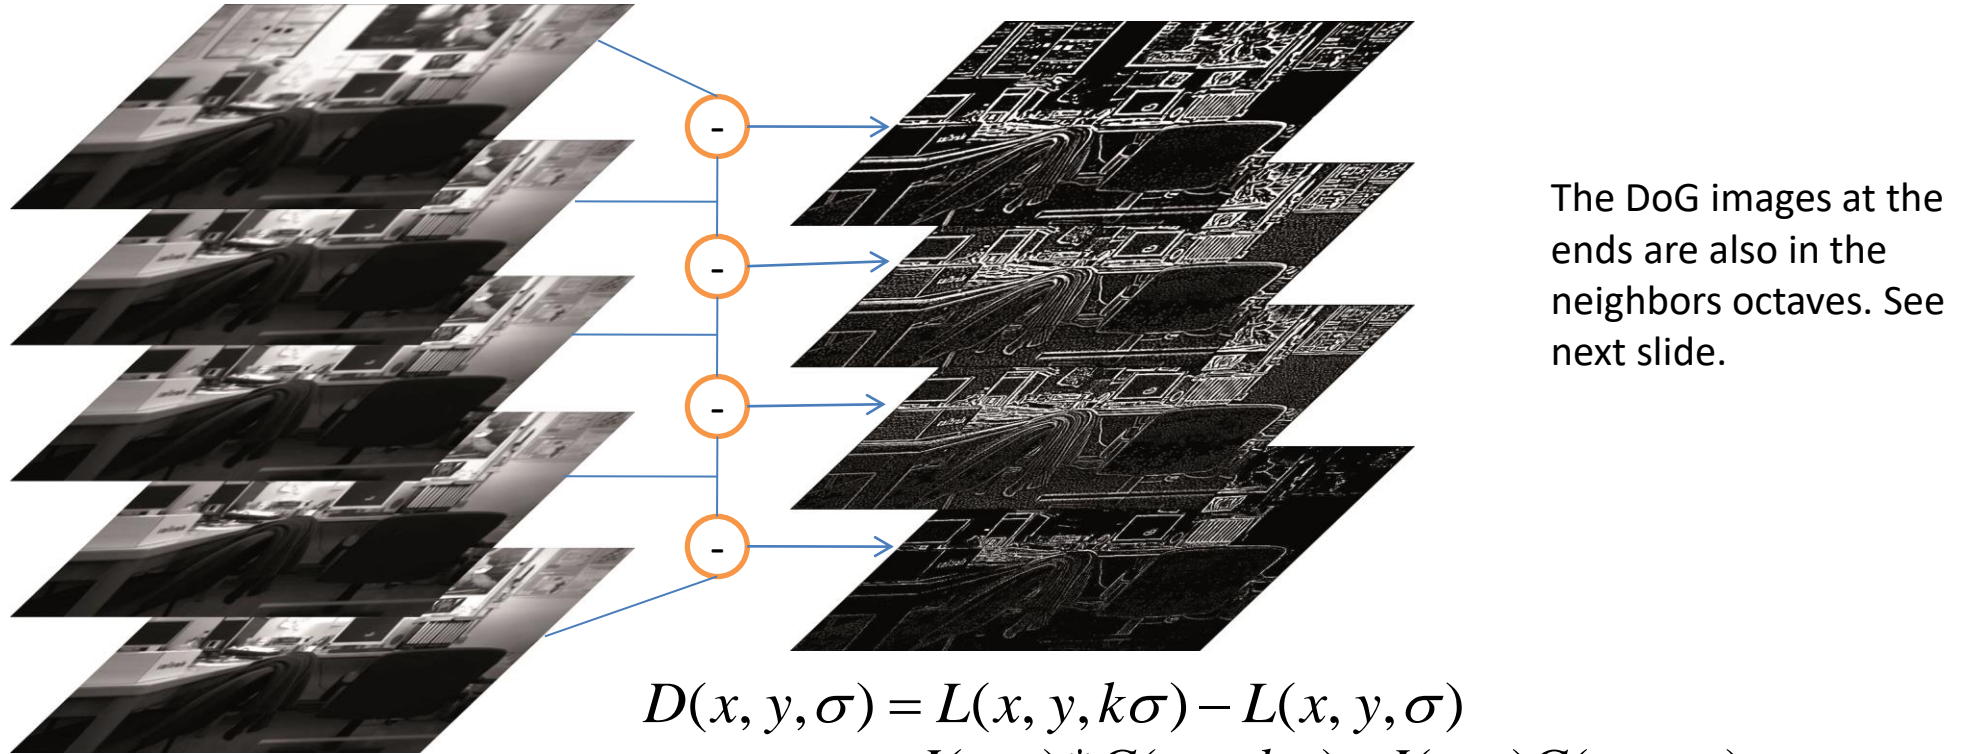
5. The SIFT operator

- The scale space has the structure of a **pyramid**: a collection of digital images sampled at progressively coarser spatial resolution and hence of progressively smaller size.
- The pyramid is built upon a number of **Octaves**.
- Each **Octave (o)** consists of $s+2$ images of the **same size (resolution)** with increasing smoothness.
- In the **following Octave** the image has half the resolution (size) since it does not make sense to keep the resolution when small details have been removed.



5. The SIFT operator

From smoothed images to **Difference of Gaussians (DoG)**




$$\begin{aligned} D(x, y, \sigma) &= L(x, y, k\sigma) - L(x, y, \sigma) \\ &= I(x, y) * G(x, y, k\sigma) - I(x, y) * G(x, y, \sigma) \\ &= I(x, y) * [G(x, y, k\sigma) - G(x, y, \sigma)] \\ &= I(x, y) * DoG(x, y, \sigma) \end{aligned}$$

Diference of smoothed images = Image convolved with a DoG

5. The SIFT operator

Construction of the scale space through “octaves”




DoG are used here to detect BLOBS (not edged)!!

5. The SIFT operator

Search for extreme points

Each pixel value is compared to its 26 neighbors along the full scale:

- **8** in the same scale,
- **9** in the upper scale and
- **9** in the lower scale




A extreme point in the scale gives us a distinctive point in (x,y) and in the pyramid

A extreme point in the scale gives us a distinctive point in (x,y) and in the pyramid (k)


WHY?

Recall:




DoG operator (1D)

2D




1D




5. The SIFT operator


Example: SIFT vs. Harris



SIFT




Harris




SIFT



SIFT



SIFT



SIFT

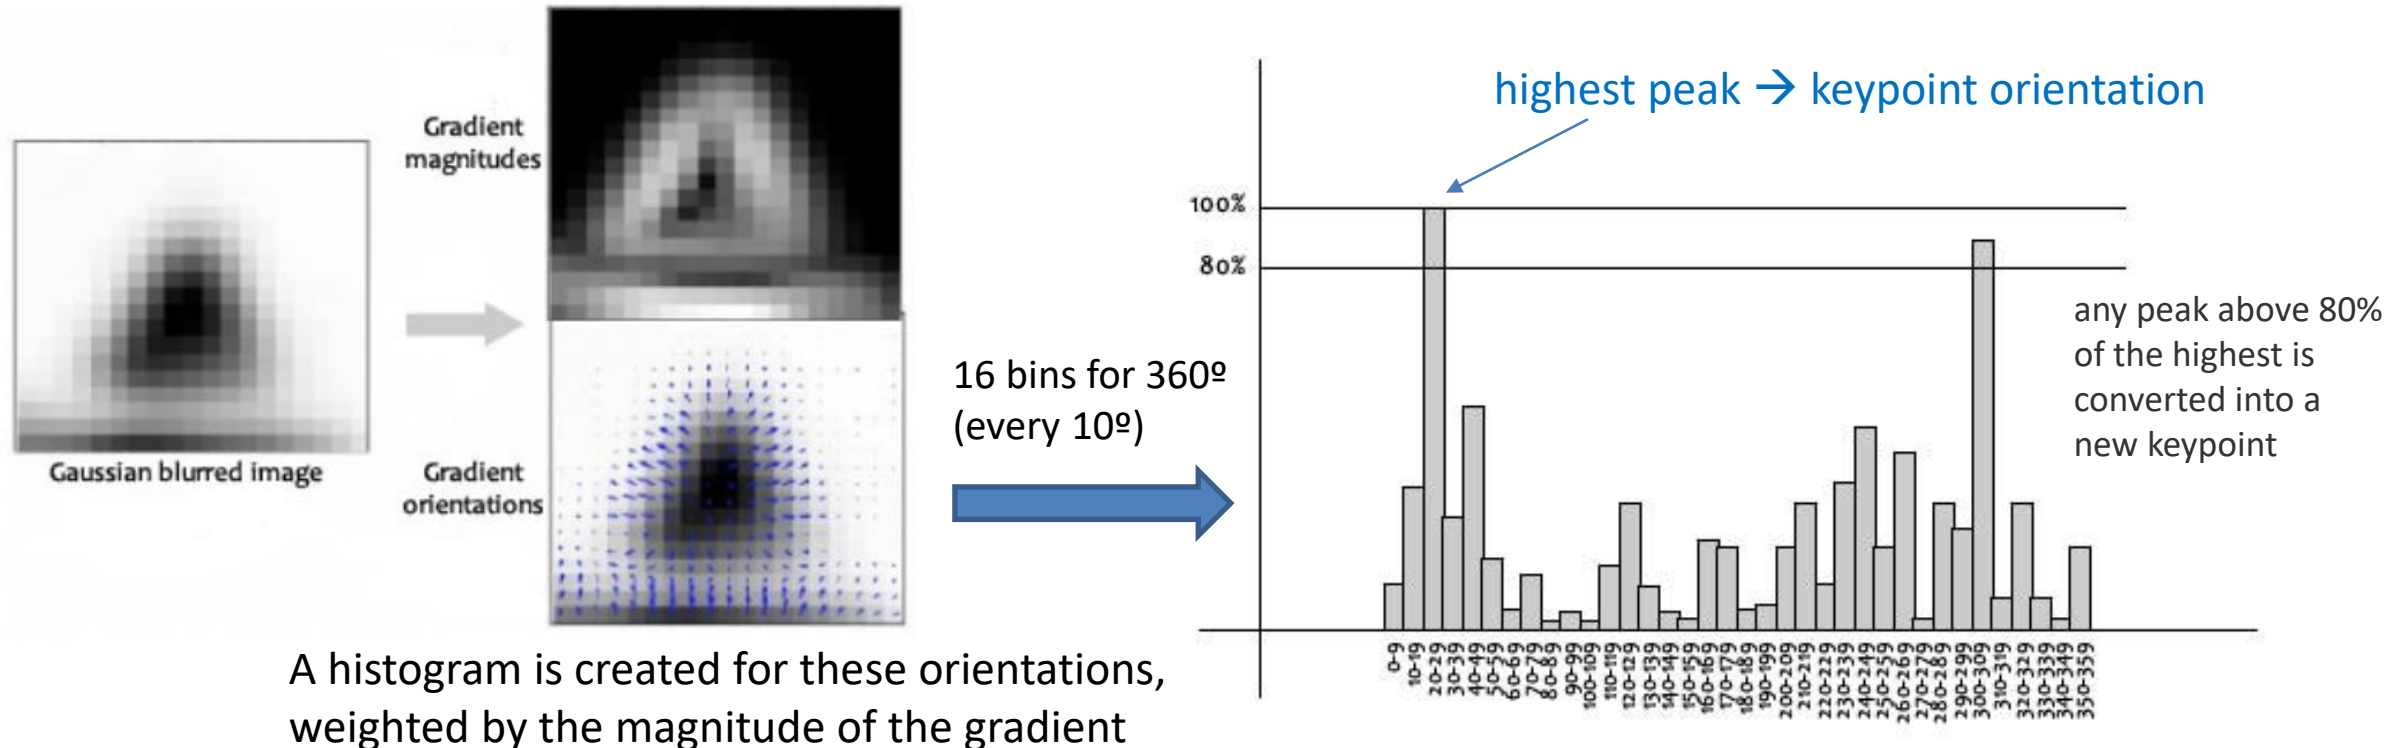
SIFT Descriptor: Histogram of orientations around the extreme point

Obtaining the **descriptor orientation**:

The magnitude (m) and orientation (θ) of the gradient is calculated for all pixels around the keypoint.

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$

$$\theta(x, y) = \tan^{-1}((L(x, y+1) - L(x, y-1)) / (L(x+1, y) - L(x-1, y)))$$




SIFT Descriptor: Histogram of orientations around the extreme point (EP)


Obtaining the **descriptor vector**:

- The neighborhood of the EP is divided in **m x m cells** (a cell is 4x4 pixels)
- Histogram of 8 orientations for each cell → Size of descriptor vector: **m x m x 8**

Example: $m = 2 \rightarrow$ Descriptor vector: $2 \times 2 \times 8 = 32$ entries




In the Lowe's original paper (D. Lowe): **m = 4** → Descriptor: $4 \times 4 \times 8 = 128D$



Obtaining the **descriptor vector**:

The histogram of orientations is weighted by

- Magnitude of the gradient: more importance to strong gradients
- Gaussian centered at the extreme point: more importance to close pixels



Descriptor: $4 \times 4 \times 8 = 128D$

SIFT Invariances

- **Scale:** Window size based on the scale at which the extreme was found
- **Orientation:** Histogram rotated along the keypoint orientation: the keypoint orientation is subtracted from each orientation of the vector



5. The SIFT operator


Example: Stereo matches from Euclidean distance

A keypoint of the left image is matched to the one in the right with the closest descriptor

Points: 965/994

Matches: 359

37.20%




5. The SIFT operator

Example: Stereo matches

Points: 245/326

Matches: 89

36.33%




5. The SIFT operator

Example: Stereo matches (with rotation)

Points: 9687/7113

Matches: 1621

22.79%




5. The SIFT operator

Example: Stereo matches (with rotation)

Points: 3105/3990

Matches: 242

7,79%



Summary

- Harris operator is ...
 - a corner detector which is combined with NCC for matching in other images
 - Based on first-order image derivatives
 - invariant to rotation (because derivatives along the eigenvectors)
 - NOT invariant to scale
 - Invariant to brightness (pixel intensities are not directly considered but derivatives)
 - Robust to noise (because of gaussian smoothing)
- KLT operator
 - Same idea as Harris but the two eigenvalues are used
- SIFT operator
 - Detect Blobs at different scales based on the DoG operator
 - Provides a descriptor with the information of the gradient around the detected keypoint at its scale